

CSCE 448/748 – Computational Photography

Sampling, Frequency, and Filtering

Nima Kalantari

Outline

- **Sampling and aliasing**
- **Frequency domain**
- **Smoothing filters**
- **Antialiasing**
- **Other filters**

Image formation involves sampling

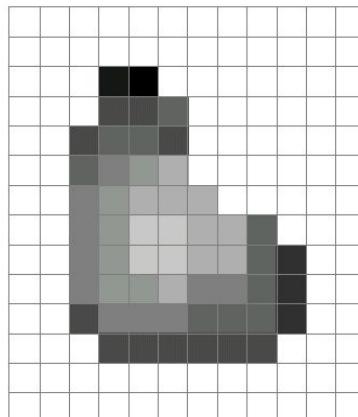
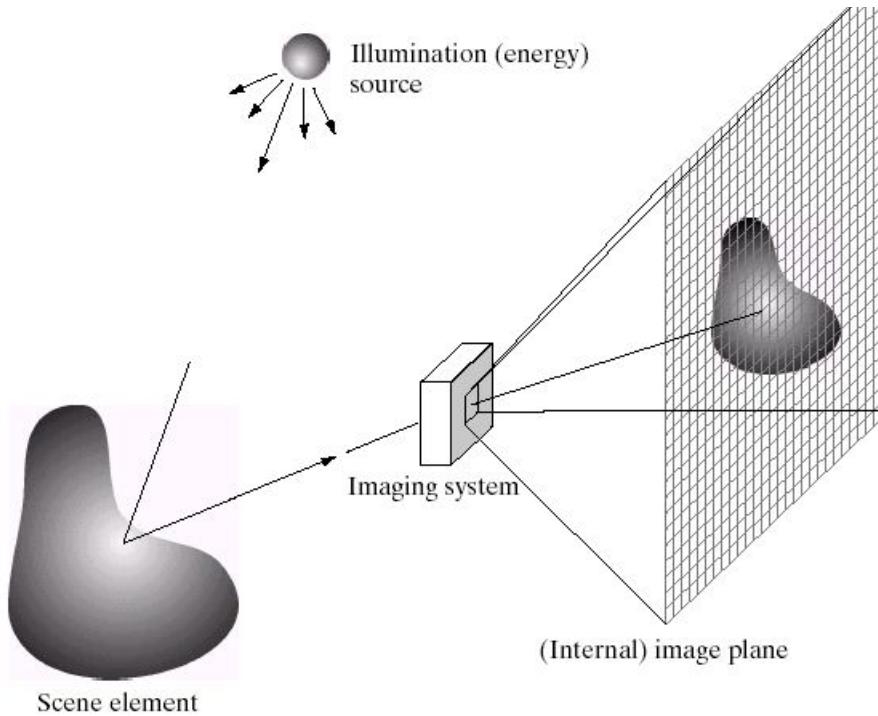
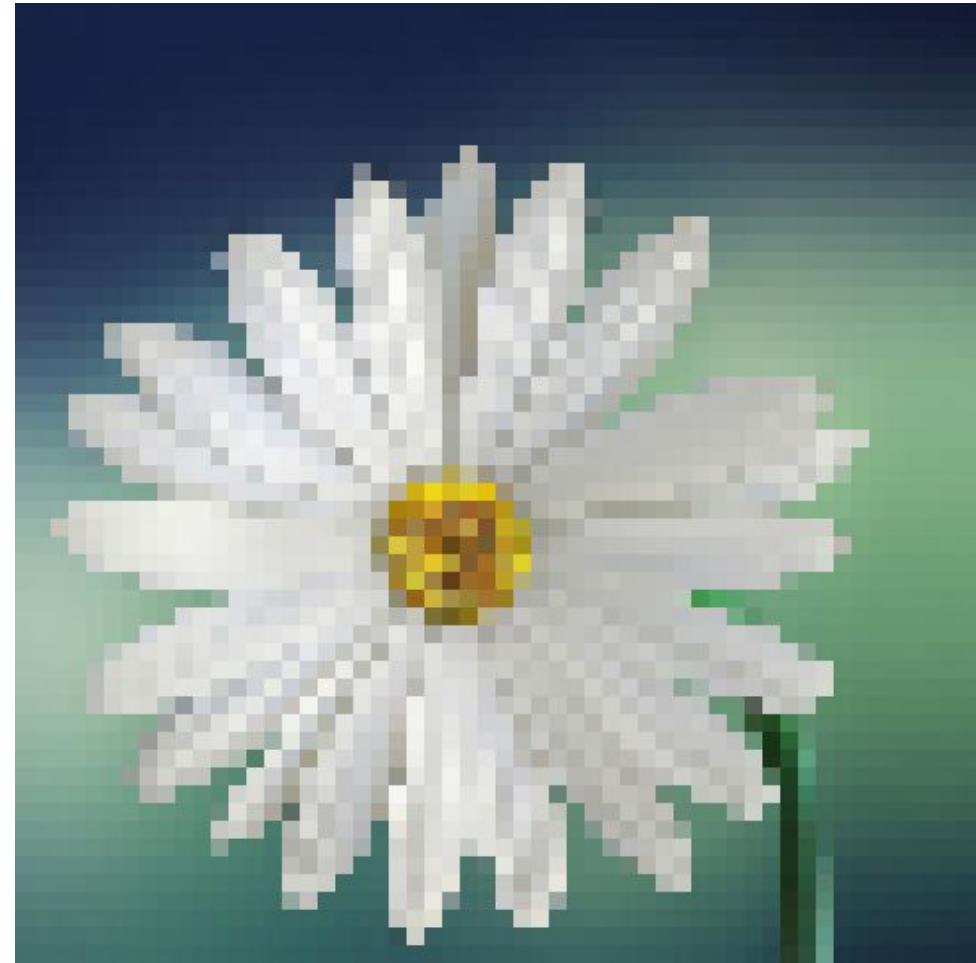


Image formation involves sampling



Scene



Digital image

Aliasing in images (Moire pattern)



Original



Aliased

Aliasing in video (wagon-wheel effect)



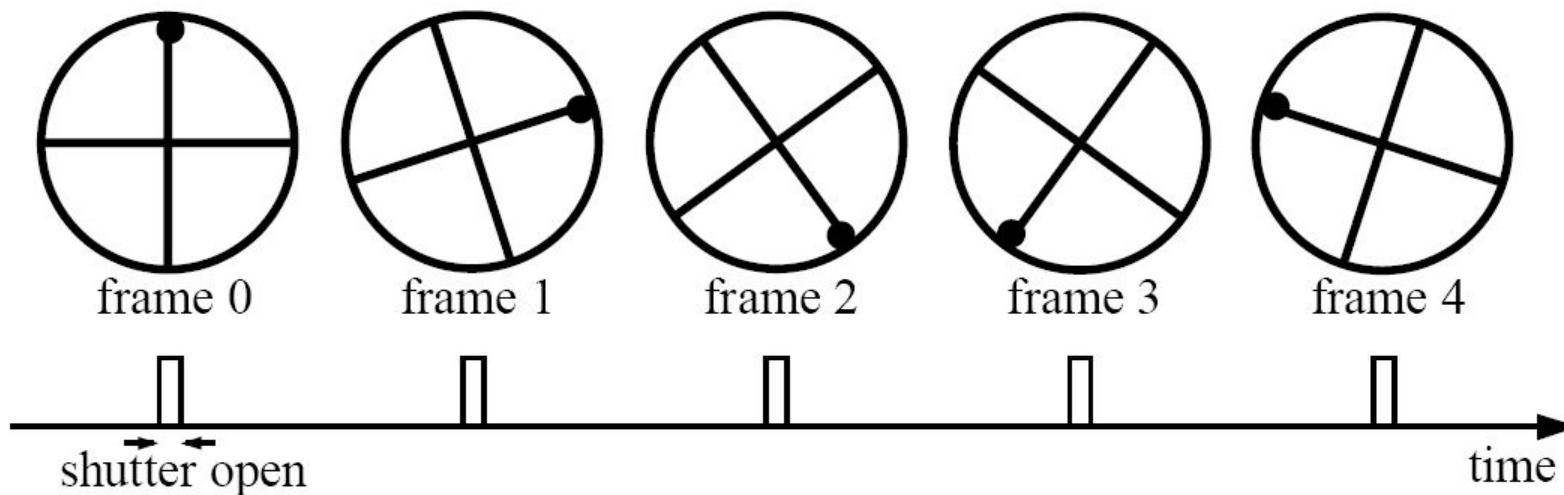
Created by Jesse Mason, https://www.youtube.com/watch?v=QOwzkND_ooU

Aliasing in video (wagon wheel effect)

Imagine a spoked wheel moving to the right (rotating clockwise).

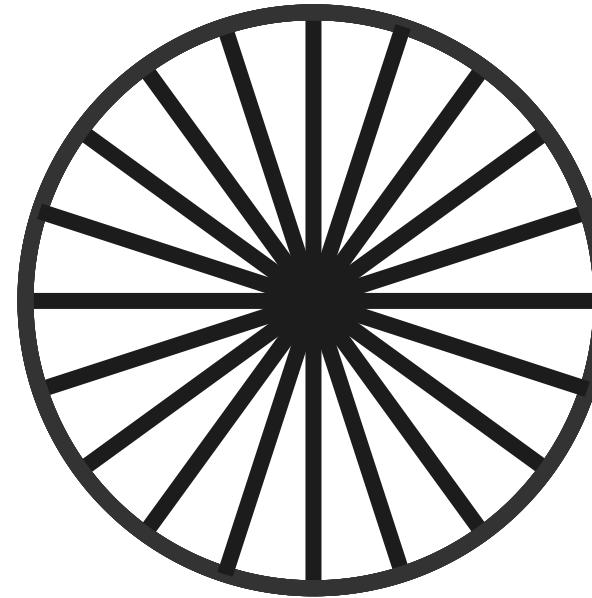
Mark wheel with dot so we can see what's happening.

If camera shutter is only open for a fraction of a frame time (frame time = 1/30 sec. for video, 1/24 sec. for film):



Without dot, wheel appears to be rotating slowly backwards!
(counterclockwise)

Aliasing in video (wagon wheel effect)



Sampling Artifacts in Computer Graphics

- **Artifacts due to sampling - “Aliasing”**
 - **Jaggies – sampling in space**
 - **Moire – undersampling images (and textures)**
 - **Wagon wheel effect – sampling in time**
 - **[Many more] ...**
- **Aliasing**
 - **Fast-changing signals (high frequency) sampled too slowly**

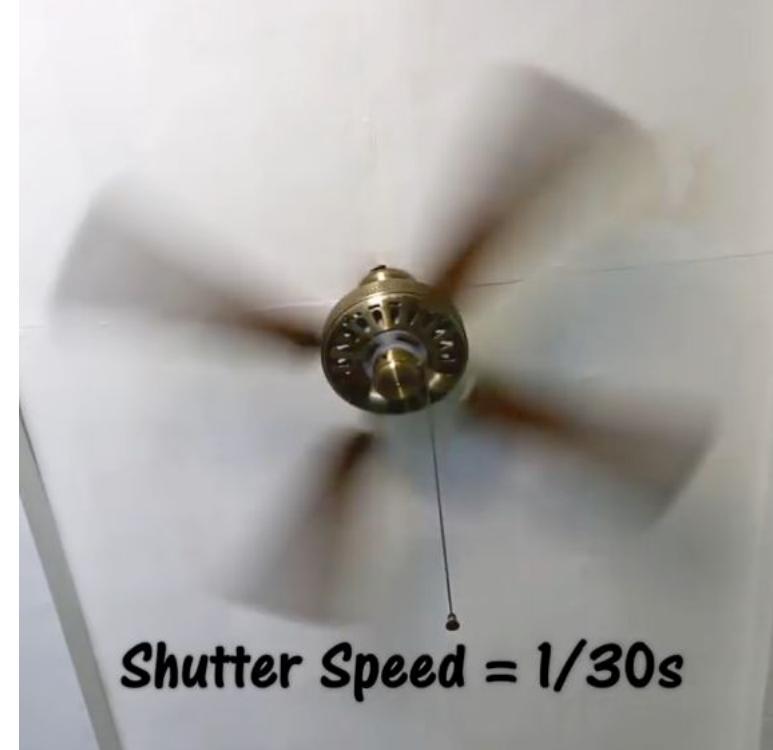
Antialiasing

Video: Point vs Antialiased Sampling



Shutter Speed = 1/800s

Point in Time



Shutter Speed = 1/30s

Motion Blurred

Video: Point Sampling in Time



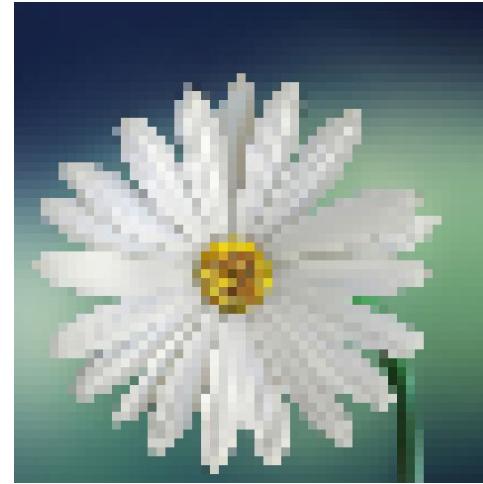
Credit: Aris & cams youtube,
<https://youtu.be/NoVwxTktoFs>

30 fps video. 1/800 second exposure is sharp in time, causes time aliasing.

Imaging: sampling in space

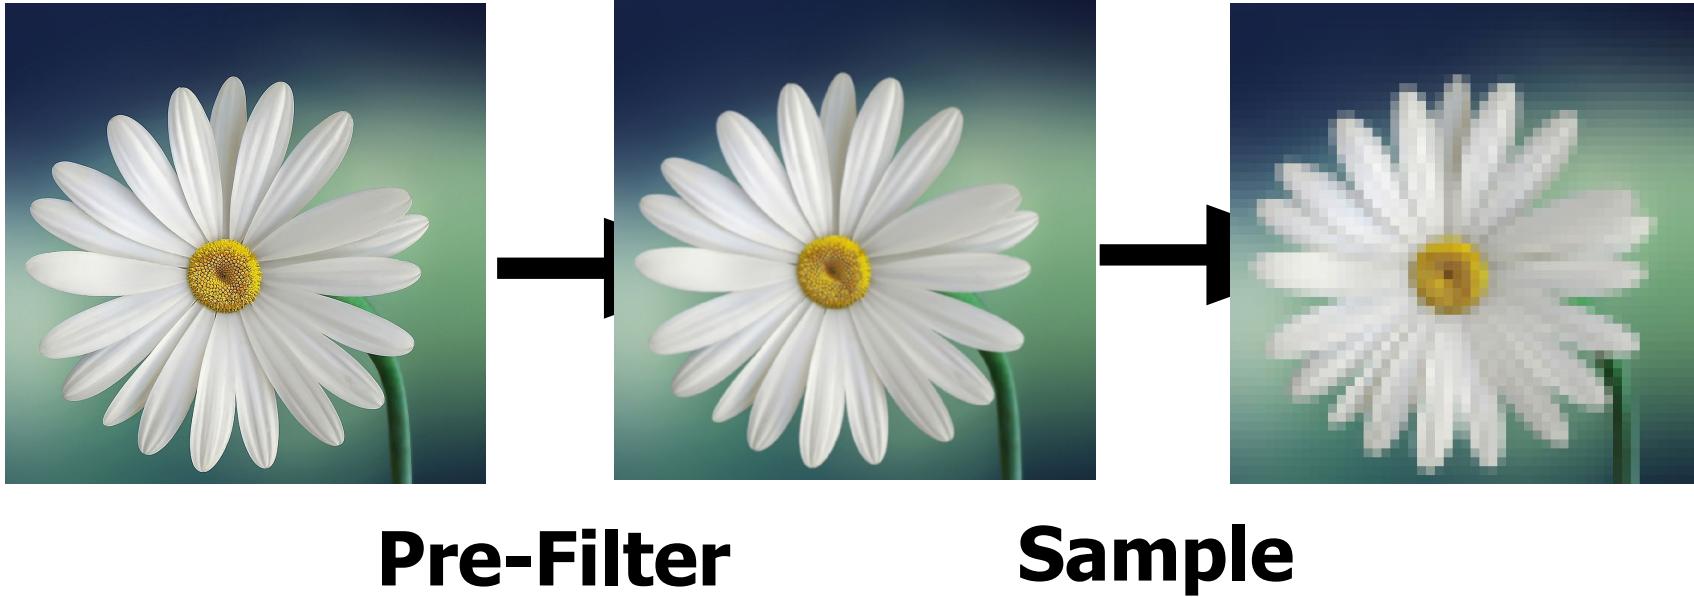


Sampl
e

A large black arrow pointing from left to right, with the word "Sampl" stacked above "e" and centered along the arrow.

Note jagged edges of the flower petals

Imaging: antialiased sampling



Aliasing and Antialiasing

- **Why undersampling results in aliasing?**
- **Why pre-filtering reduces aliasing?**

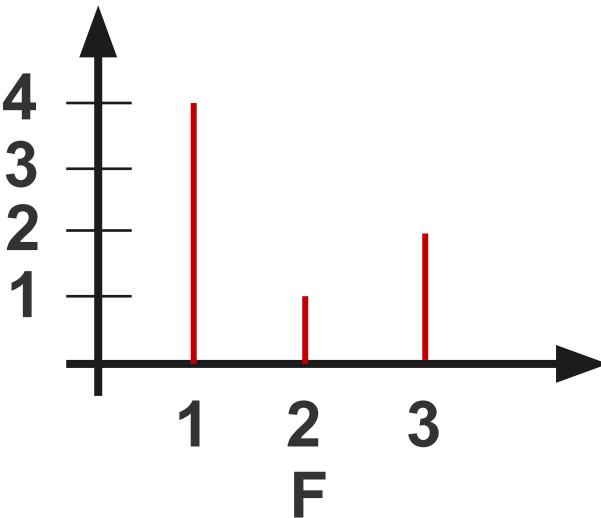
Outline

- **Sampling and aliasing**
- **Frequency domain**
- **Smoothing filters**
- **Antialiasing**
- **Other filters**

Image representation with basis functions

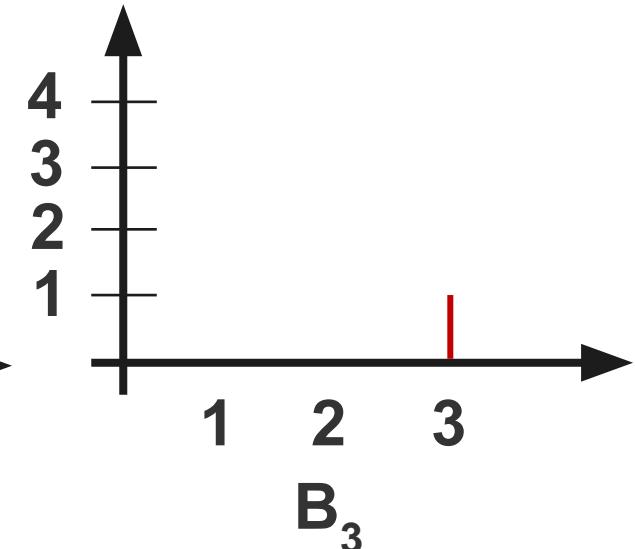
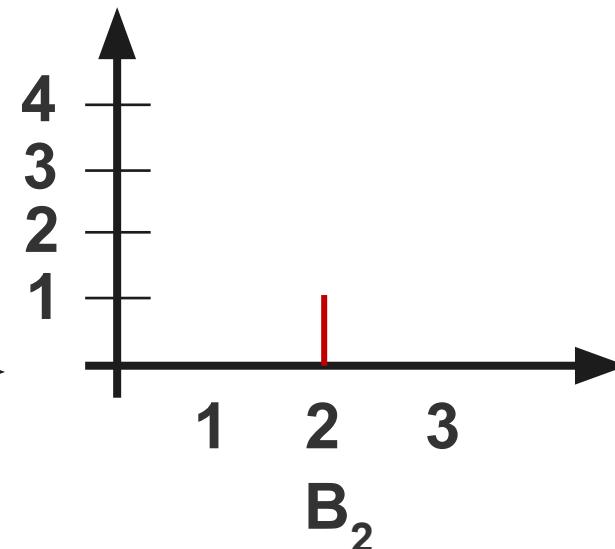
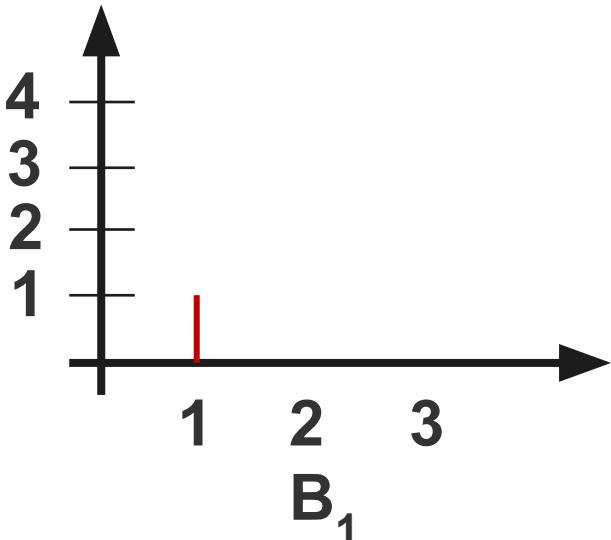
Spatial basis (1D example)

□ Signal



$$F = 4 B_1 + 1 B_2 + 2 B_3$$

□ Basis



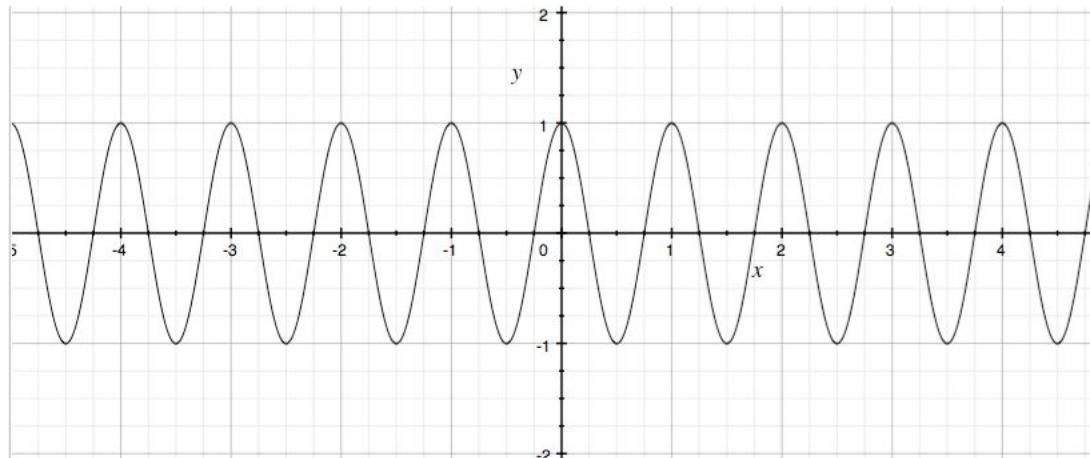
Spatial basis

- **In 2D, the basis functions will be 2 dimensional**
- **Each basis will have a value of 1 at each pixel**
- **The image can be described as a weighted combination of the basis functions**

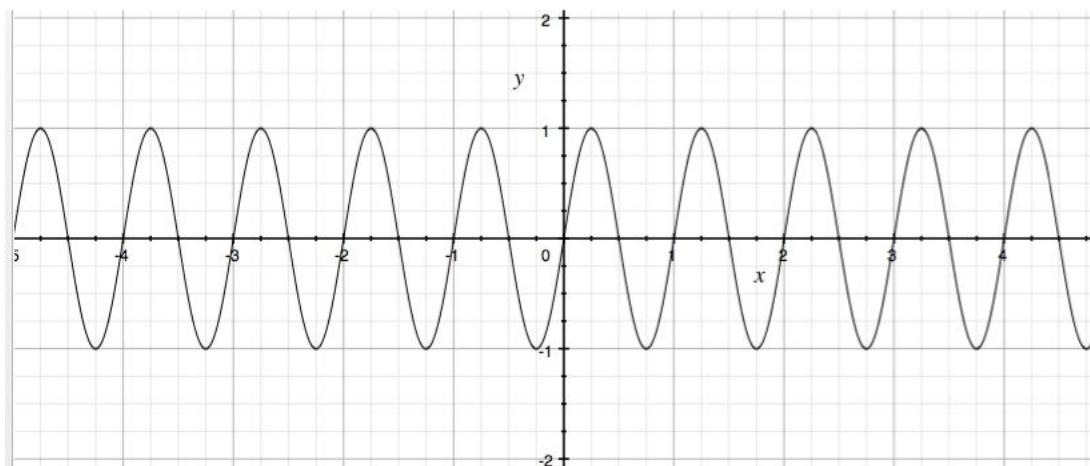
Frequency domain

- **Representing a signal using frequency basis functions**
- **The signal is represented as a weighted combination of a series of basis functions at different frequencies**

Sines and Cosines



$$\cos 2\pi x$$

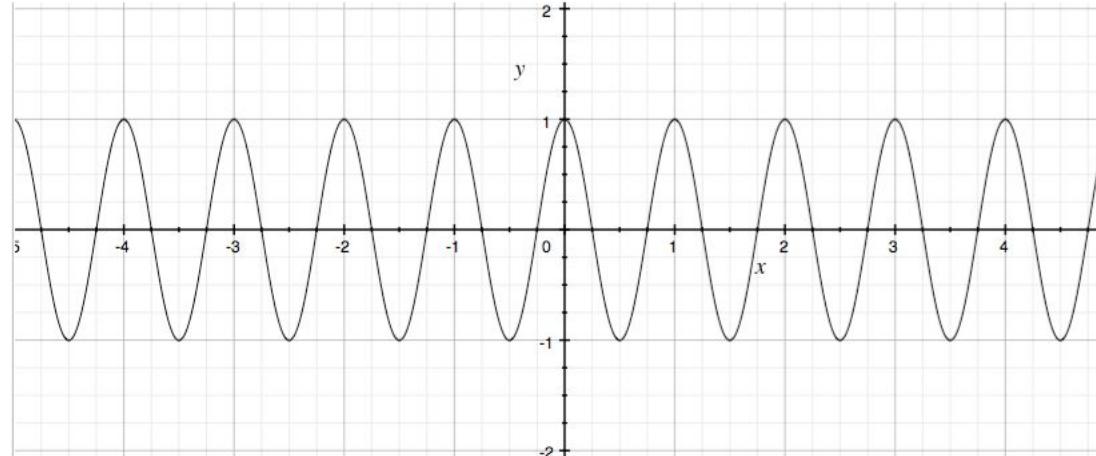


$$\sin 2\pi x$$

Frequencies

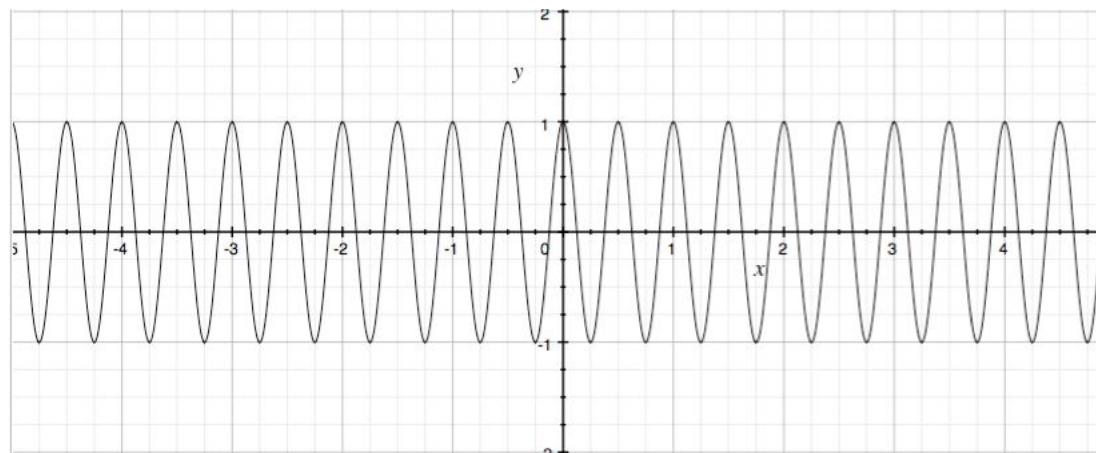
$$\cos 2\pi\omega x$$

$$\omega = \frac{1}{T}$$



$$\omega = 1$$

$$\cos 2\pi x$$



$$\omega = 2$$

$$\cos 4\pi x$$

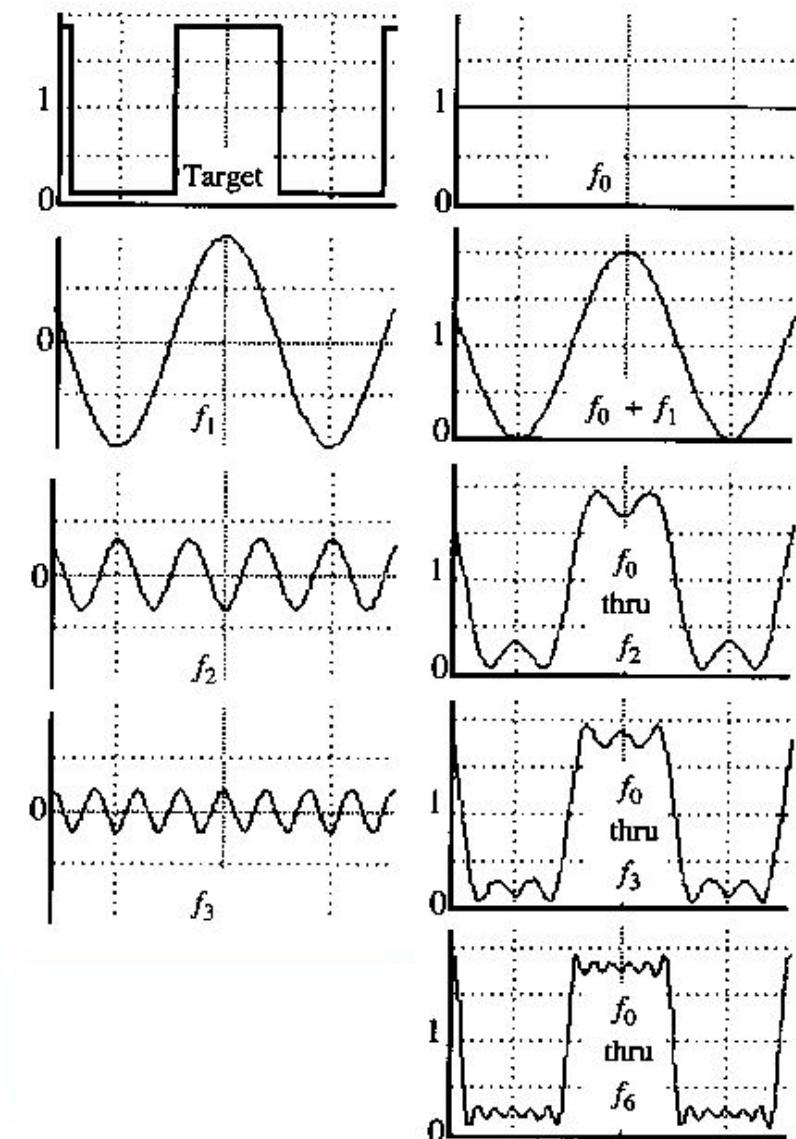
Fourier Transform

- Represent a function as a weighted sum of sines and cosines



Joseph Fourier 1768 - 1830

$$f(\text{target}) = A_0 f_0 + A_1 f_1 + A_2 f_2 + \dots$$



Extension to 2D

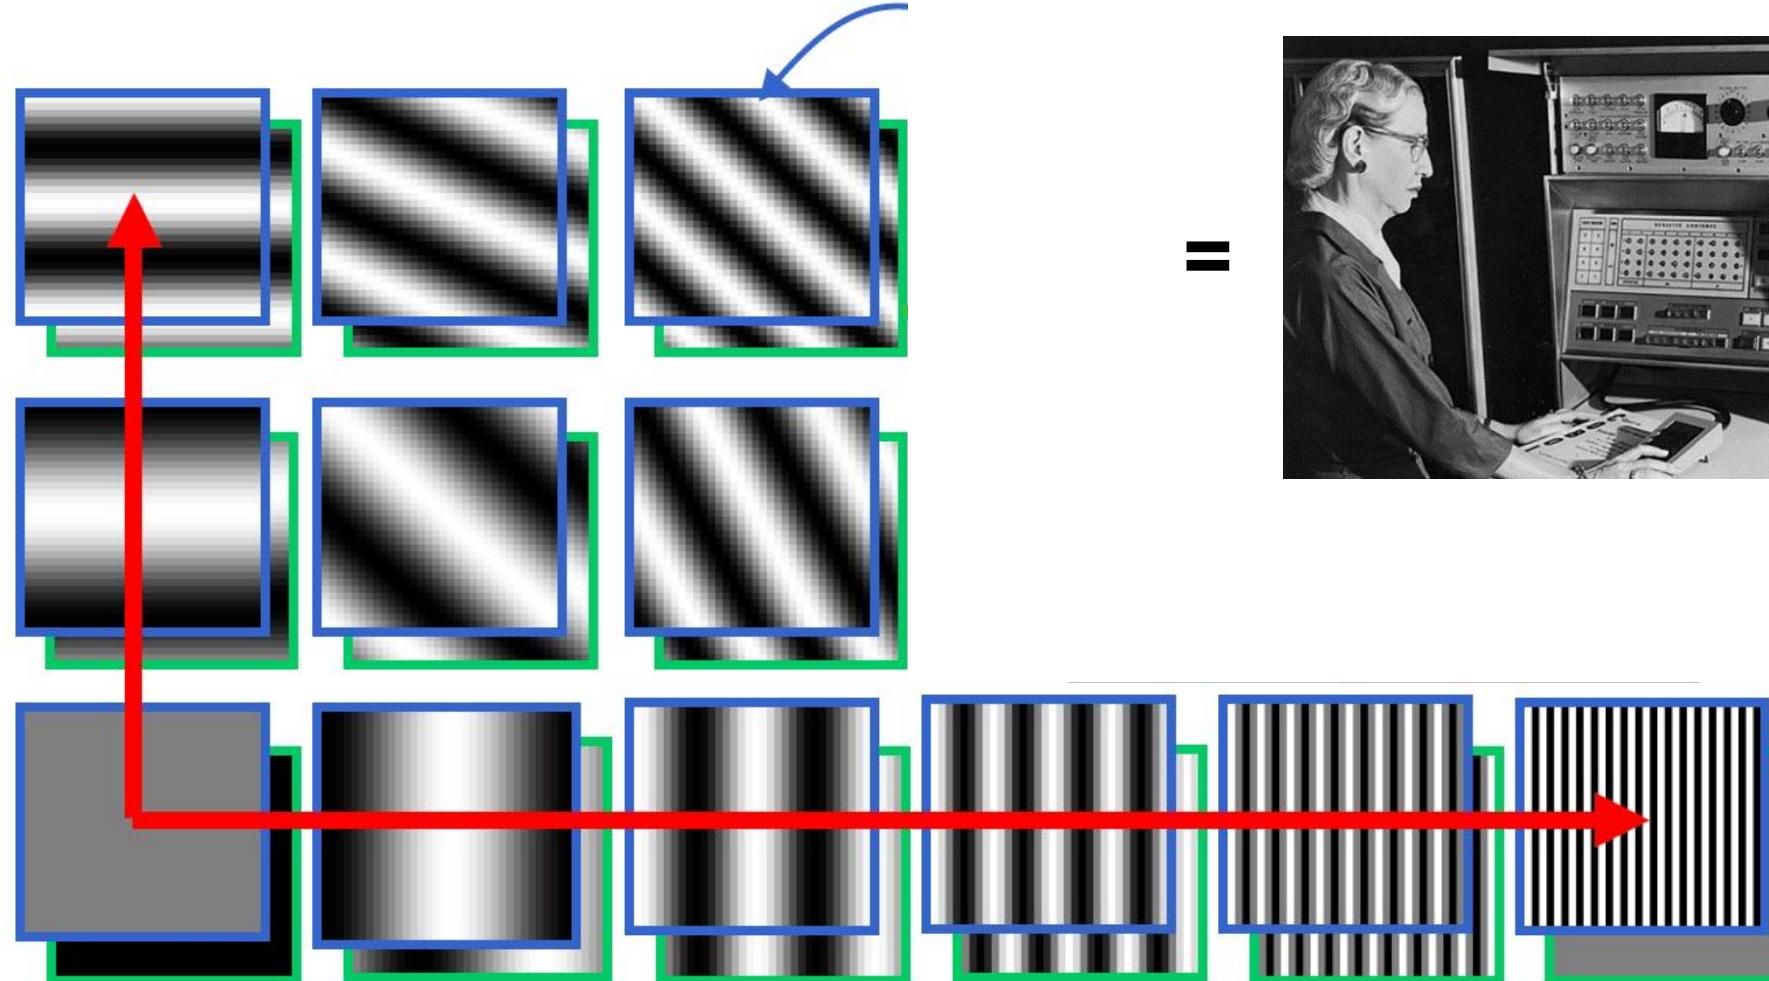
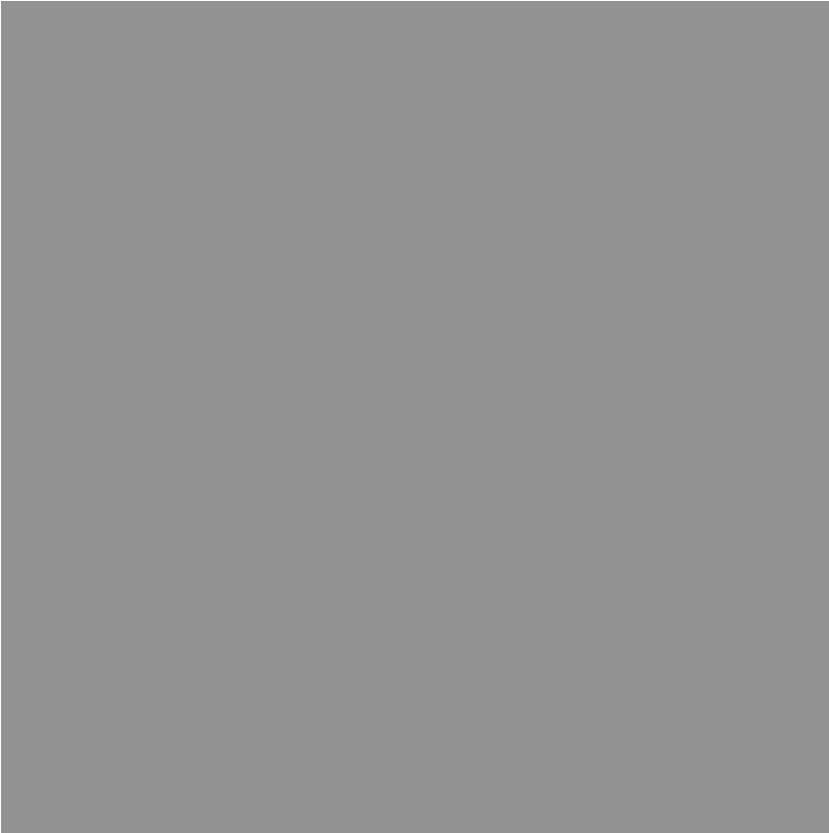
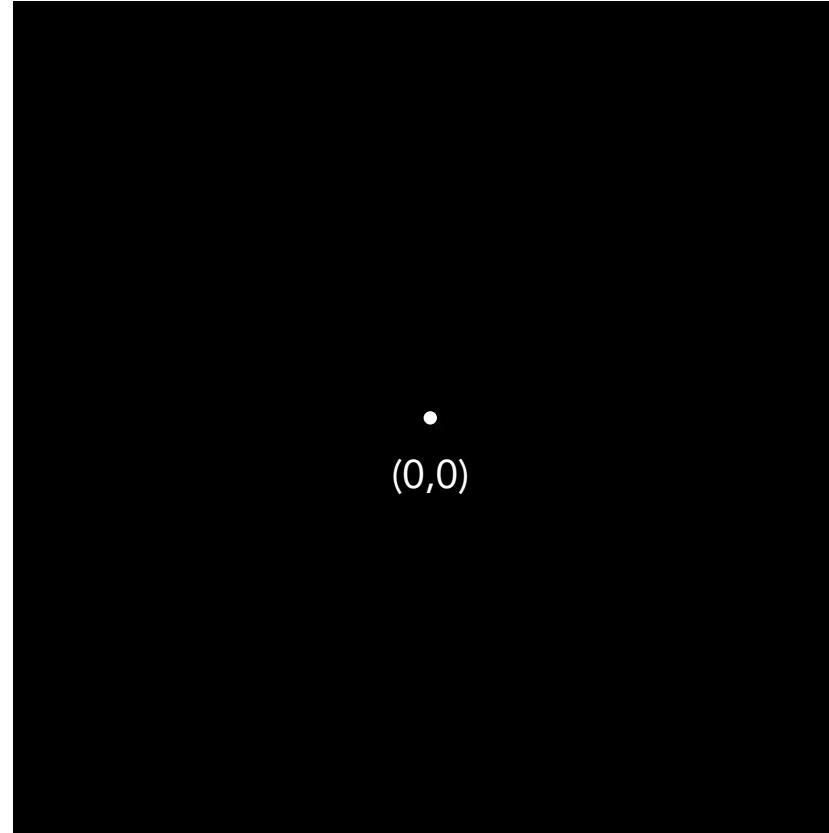


Image as a sum of basis images

Constant

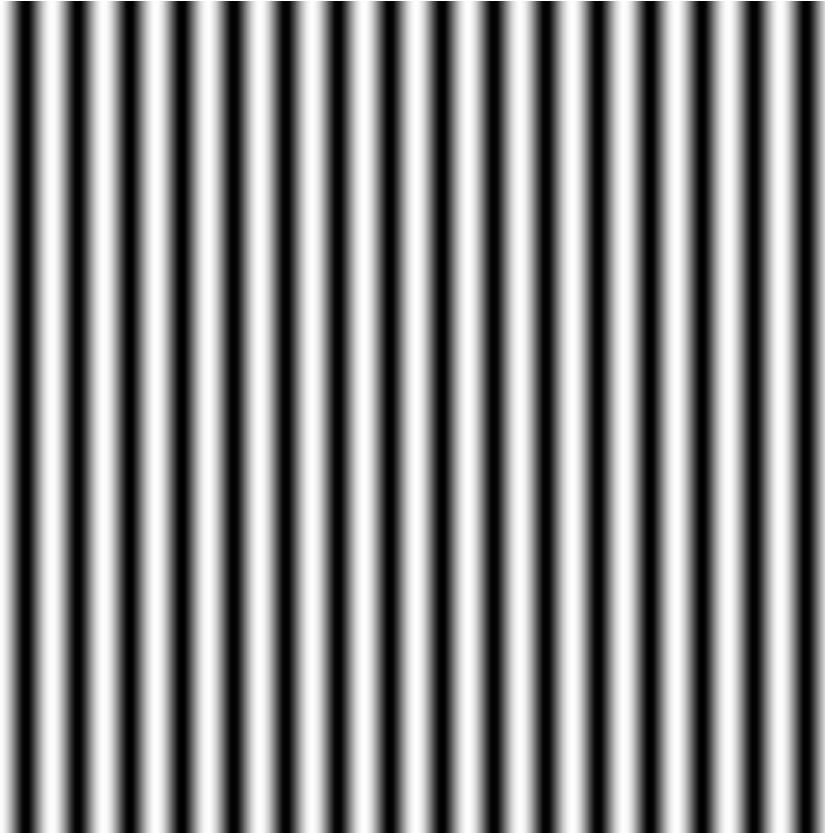


Spatial Domain

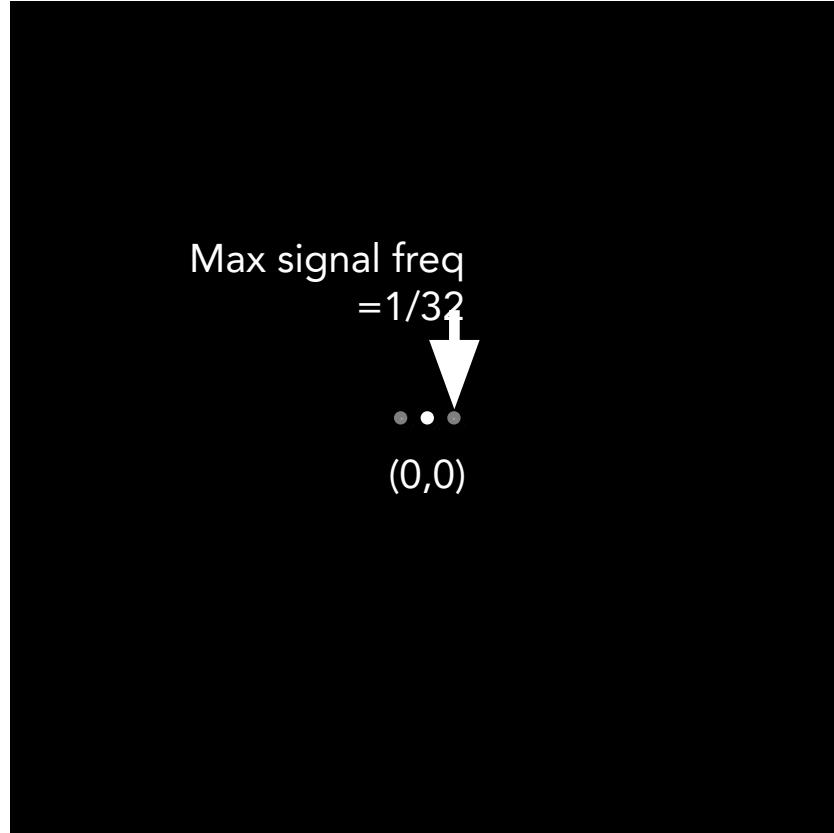


Frequency Domain

$\sin(2\pi/32)x$ — frequency 1/32; 32 pixels per cycle

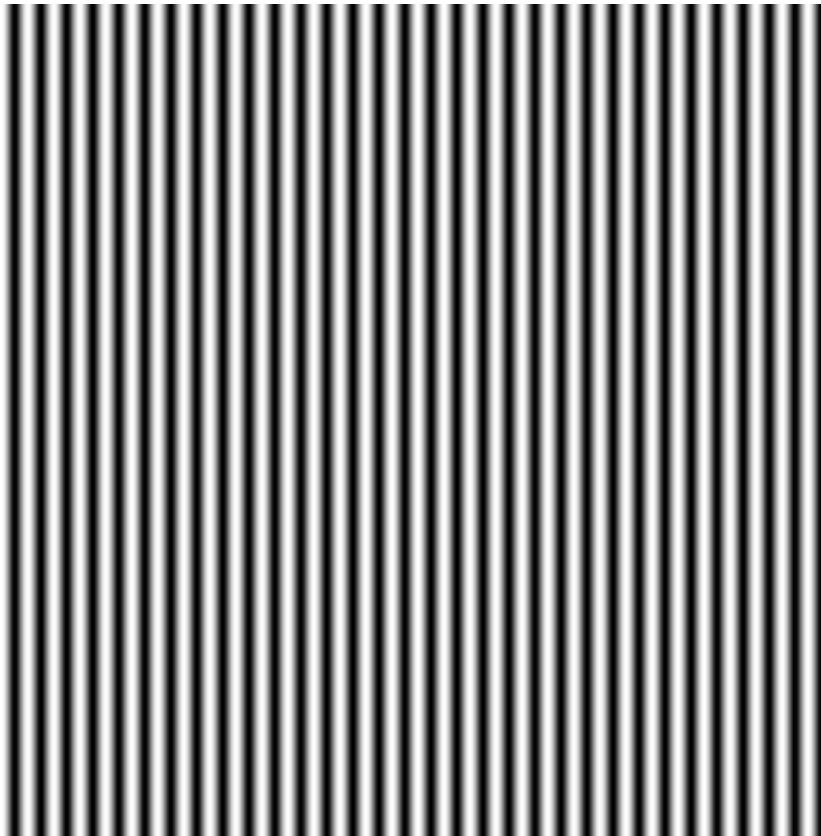


Spatial Domain

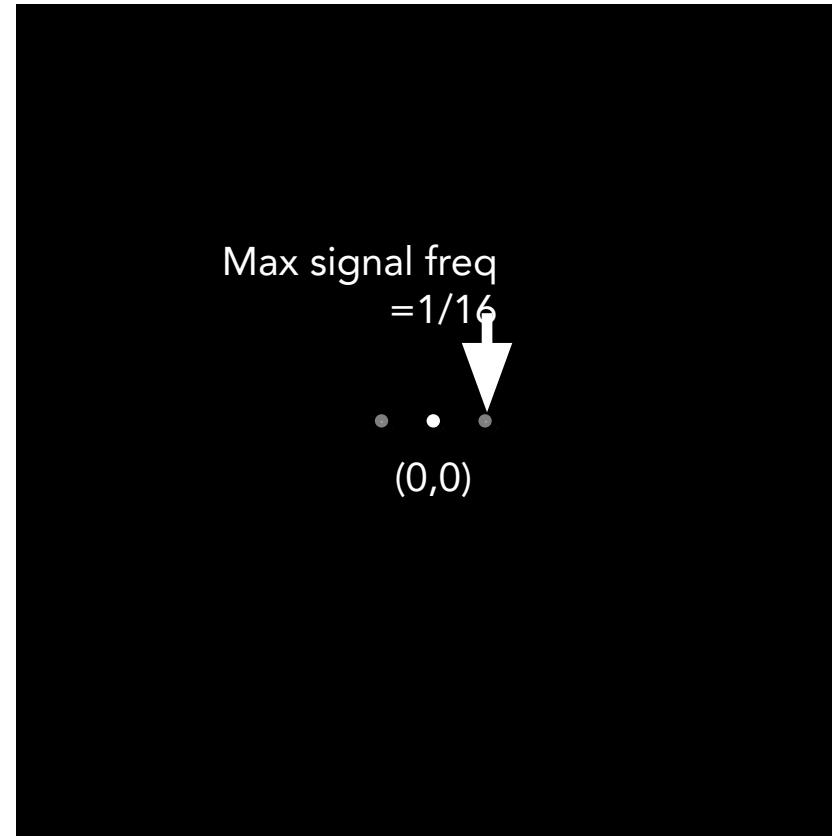


Frequency Domain

$\sin(2\pi/16)x$ — frequency 1/16; 16 pixels per cycle

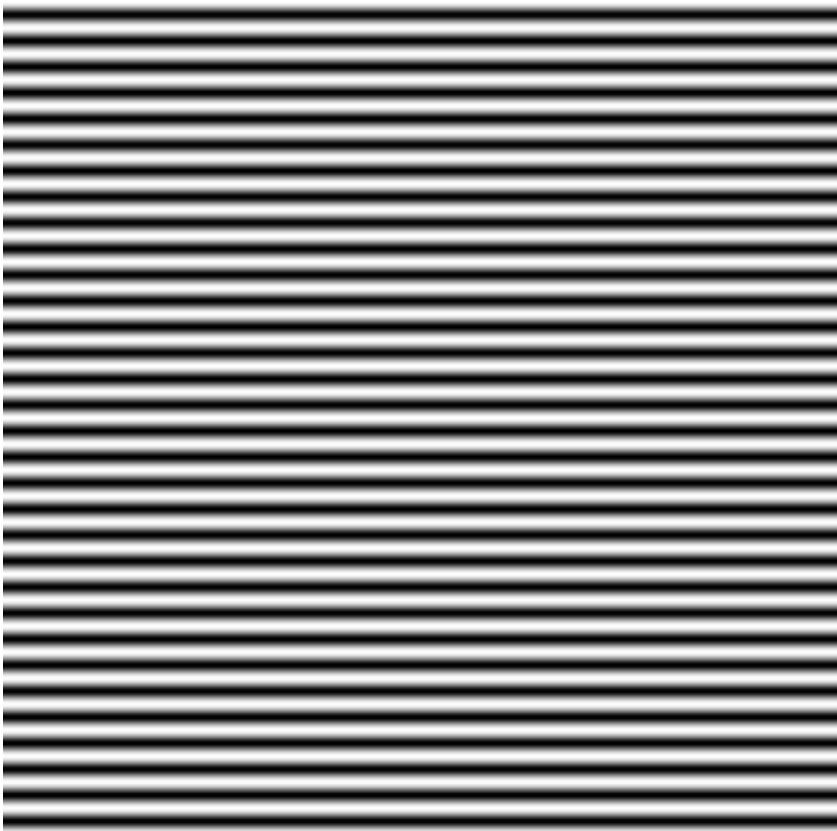


Spatial Domain

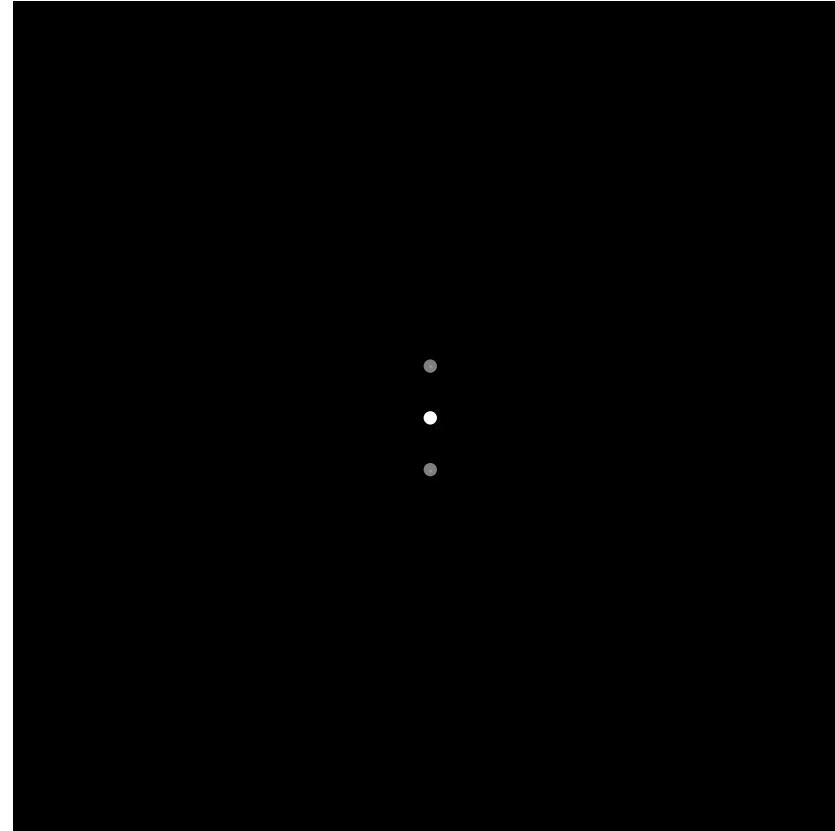


Frequency Domain

$$\sin(2\pi/16)y$$



Spatial Domain

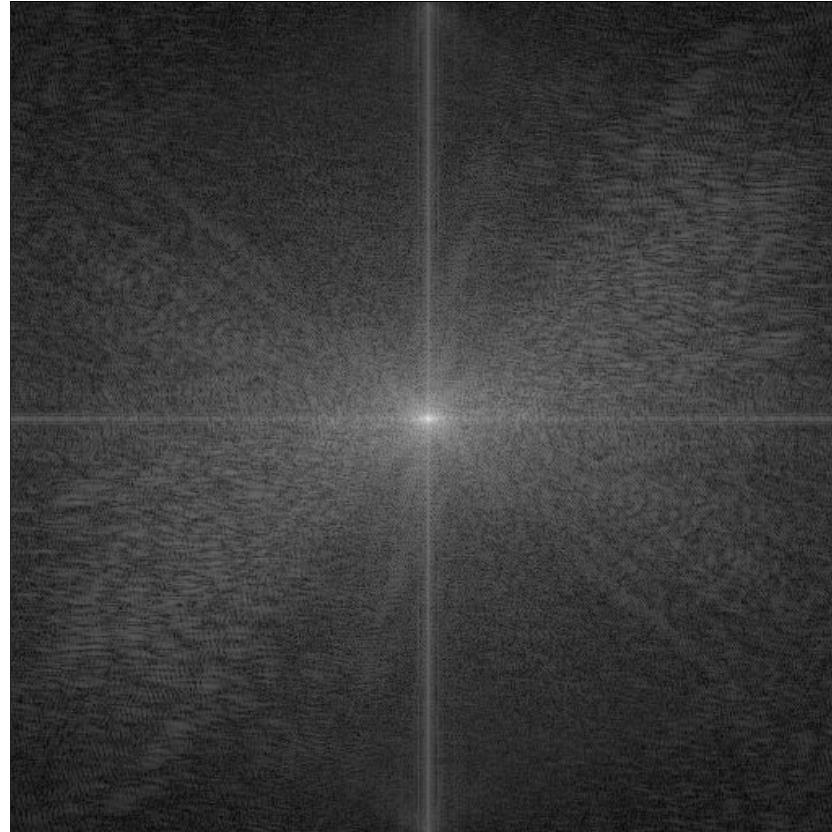


Frequency Domain

2D Frequency Space



Spatial Domain

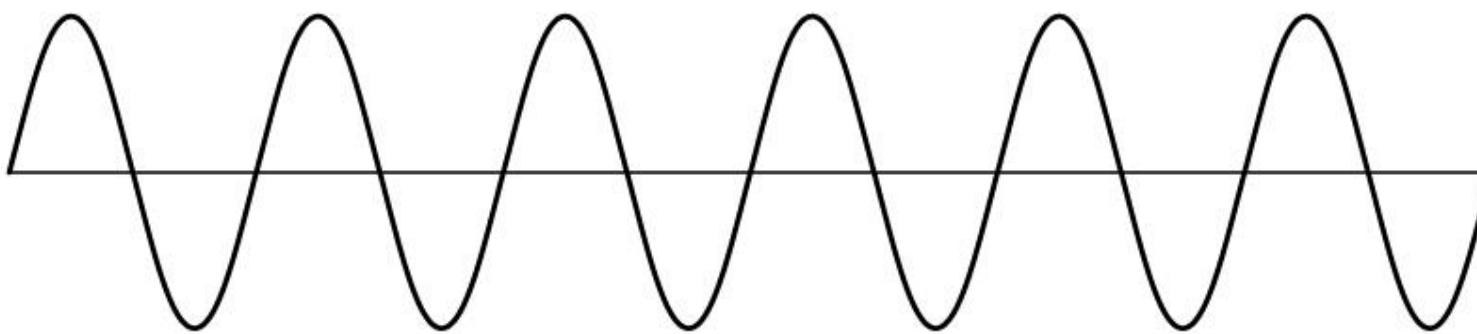


Frequency Domain

Note: Frequency domain also known as frequency space, Fourier domain, spectrum, ...

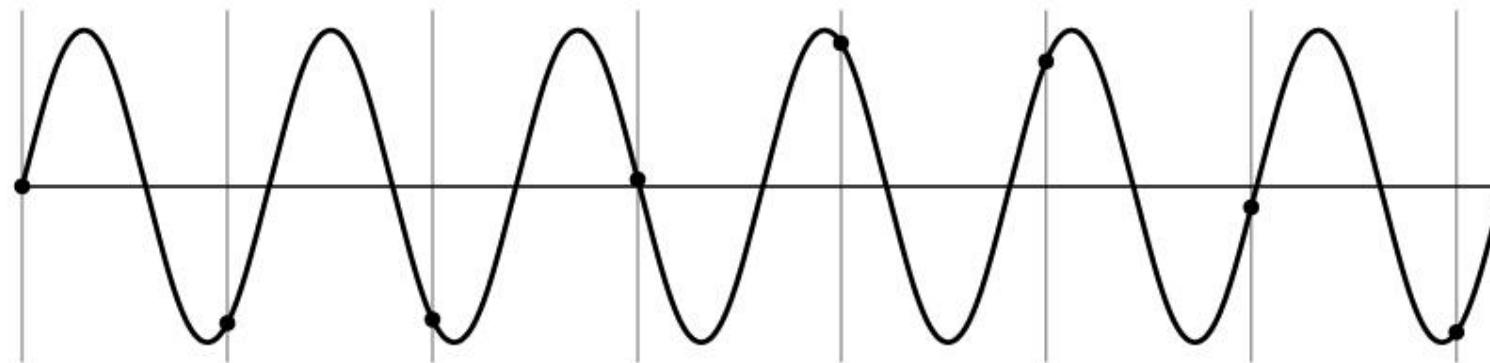
Sampling

- Simple example: a sine wave



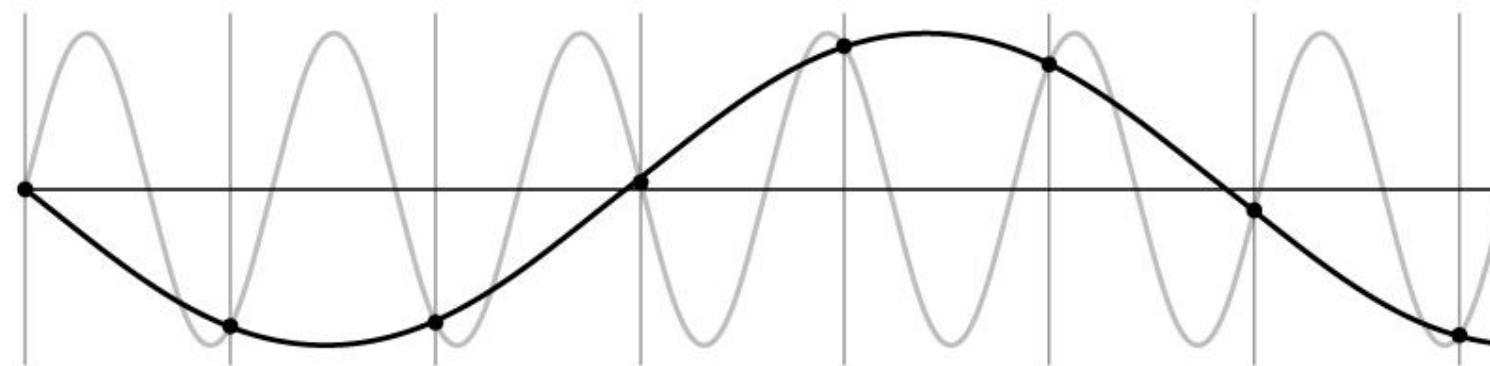
Undersampling

- What if we “missed” things between the samples?
- Simple example: undersampling a sine wave
 - Unsurprising result: information is lost



Undersampling

- What if we “missed” things between the samples?
- Simple example: undersampling a sine wave
 - Unsurprising result: information is lost
 - Surprising result: indistinguishable from lower frequency
 - Aliasing: high frequencies traveling in disguise as low frequencies



Aliasing in images



Original

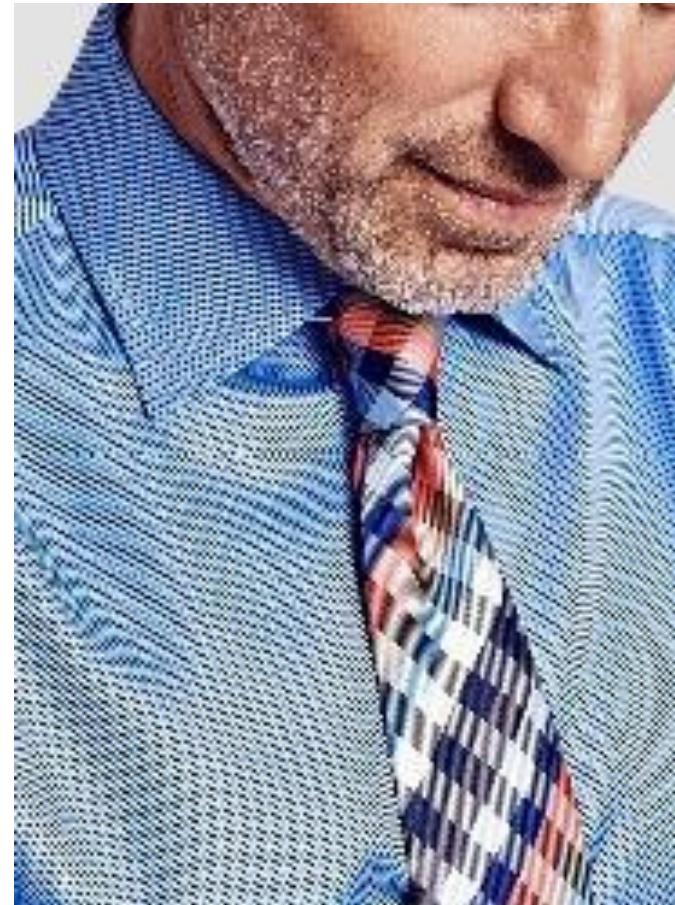


Aliased

Moiré Patterns in Imaging



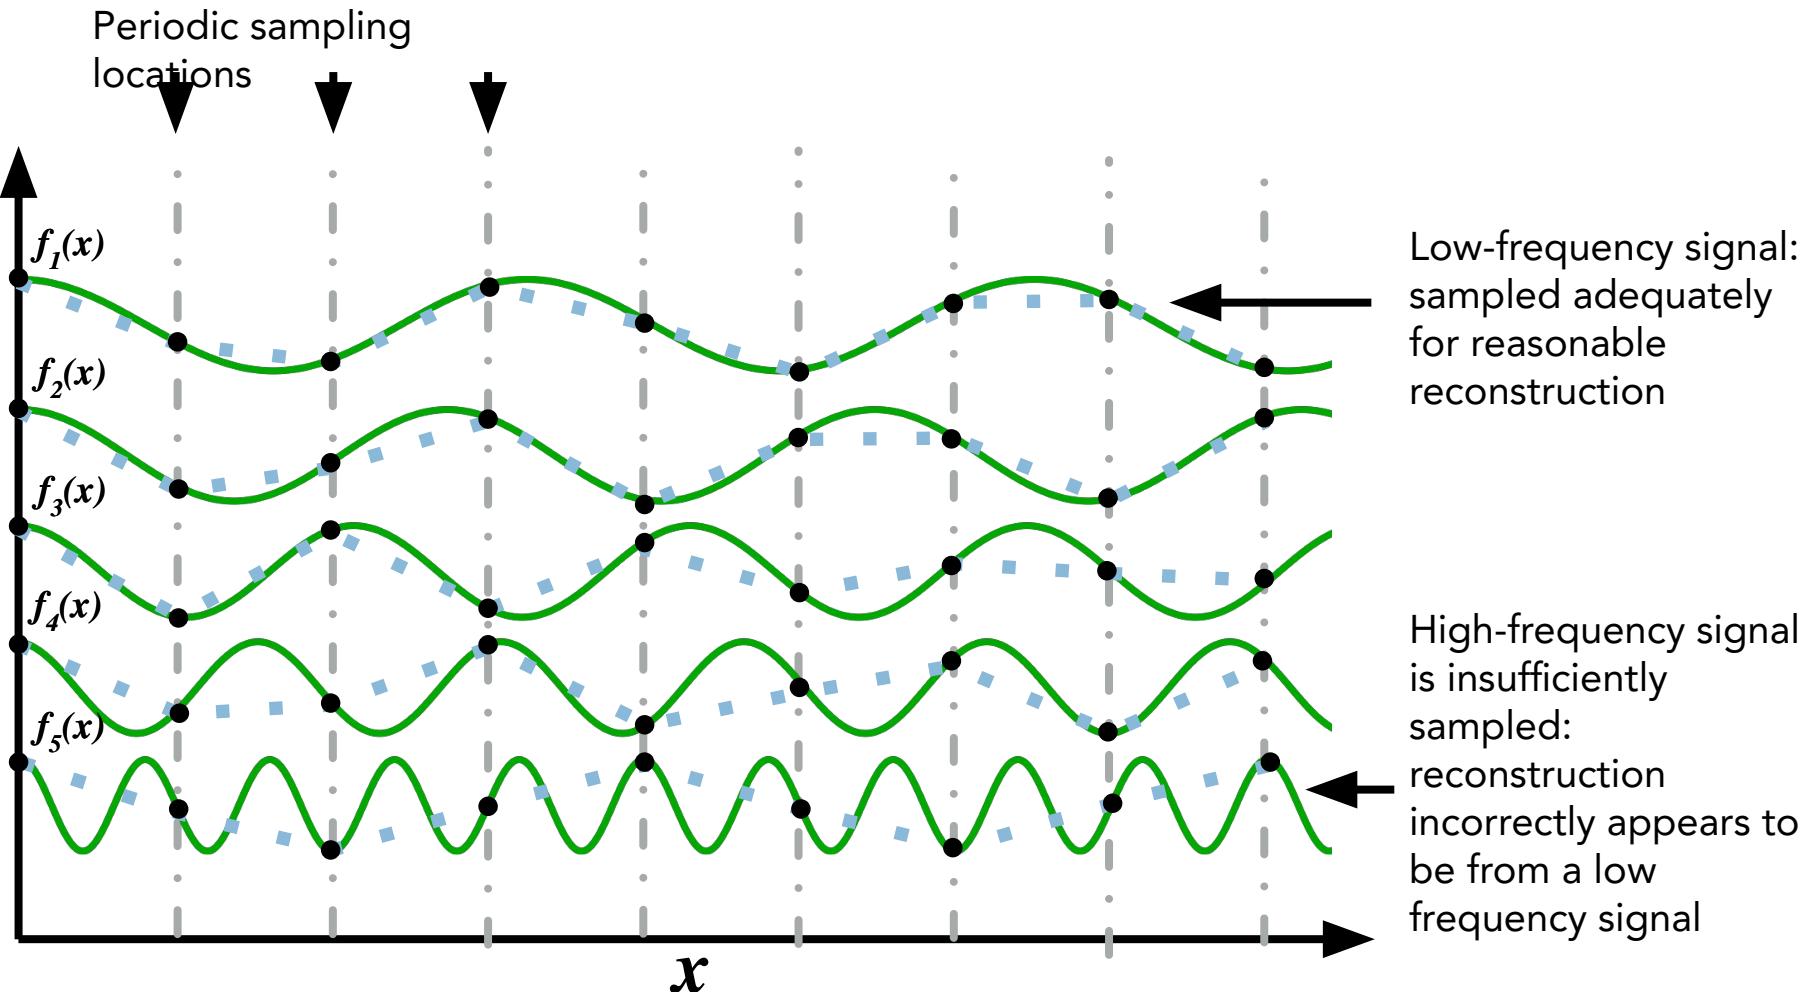
Read every sensor
pixel



Skip odd rows and
columns

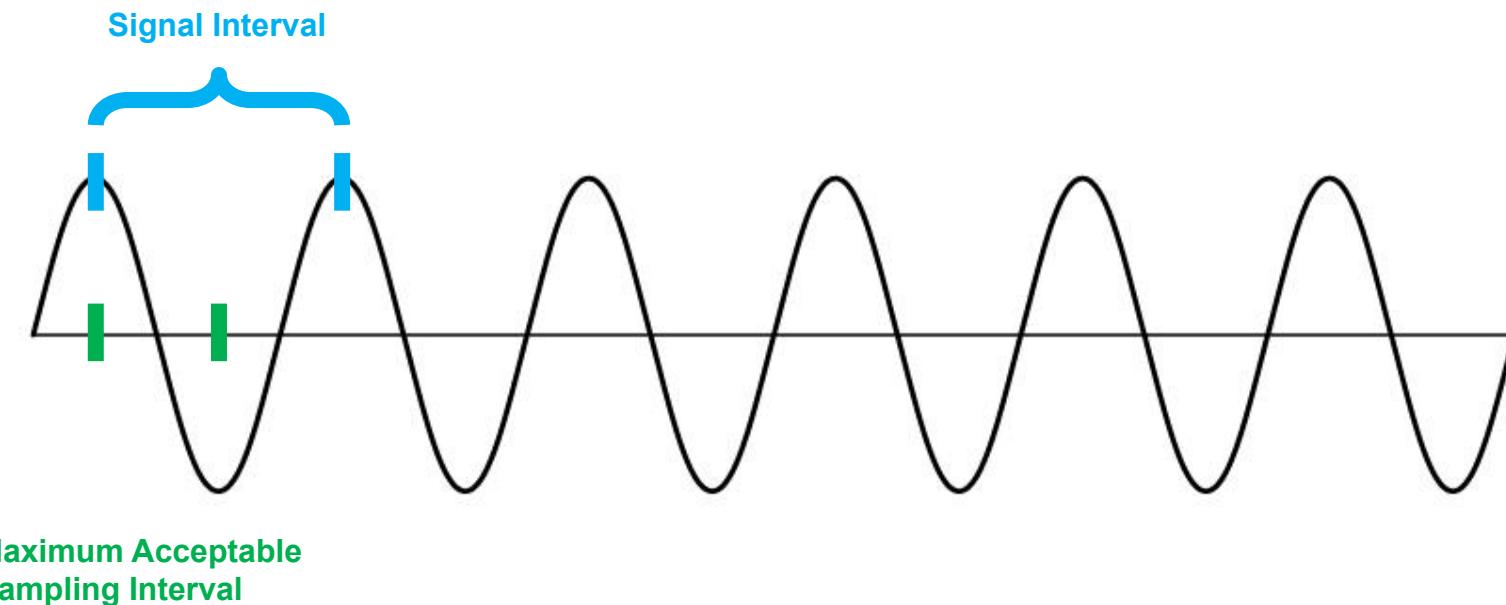
lystit.co

Higher Frequencies Need Faster Sampling



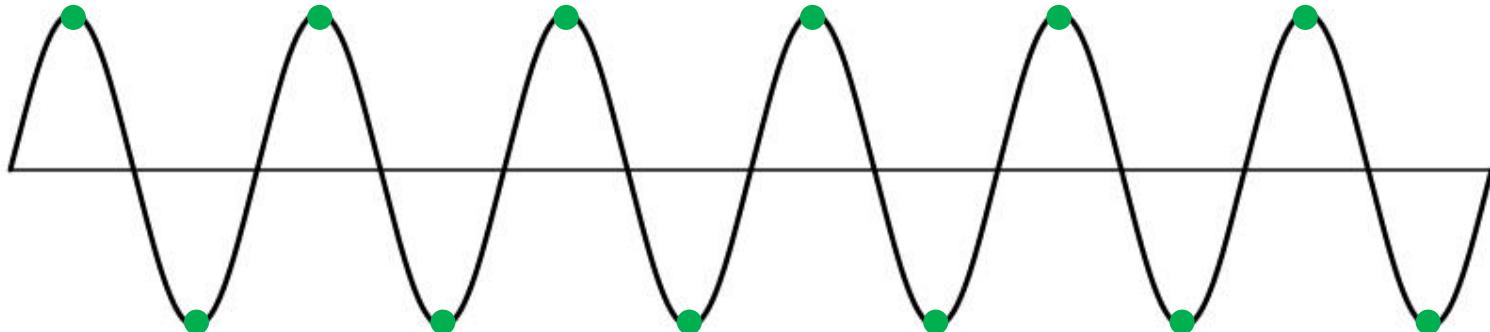
Nyquist Theorem

- Sampling rate should be equal to or greater than **twice the highest frequency** in the analog signal



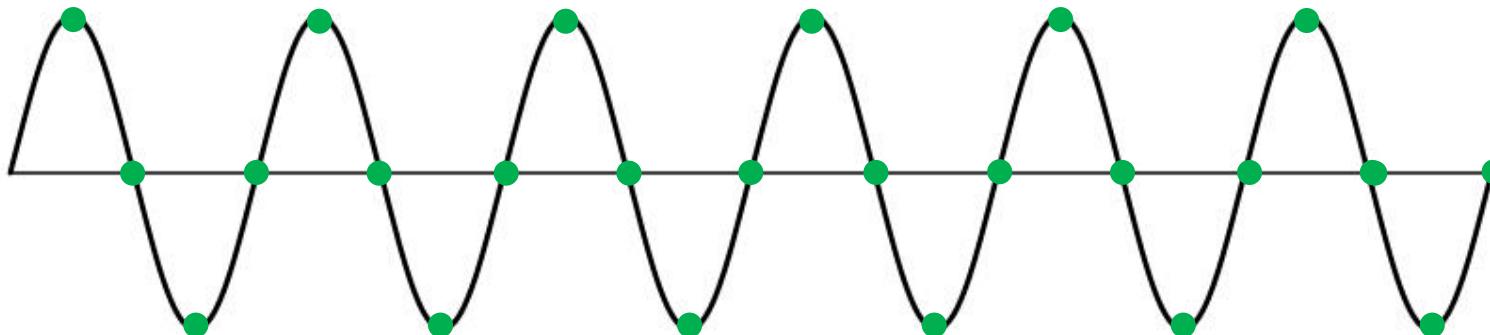
Nyquist Theorem

- Sampling rate should be equal to or greater than **twice the highest frequency** in the analog signal



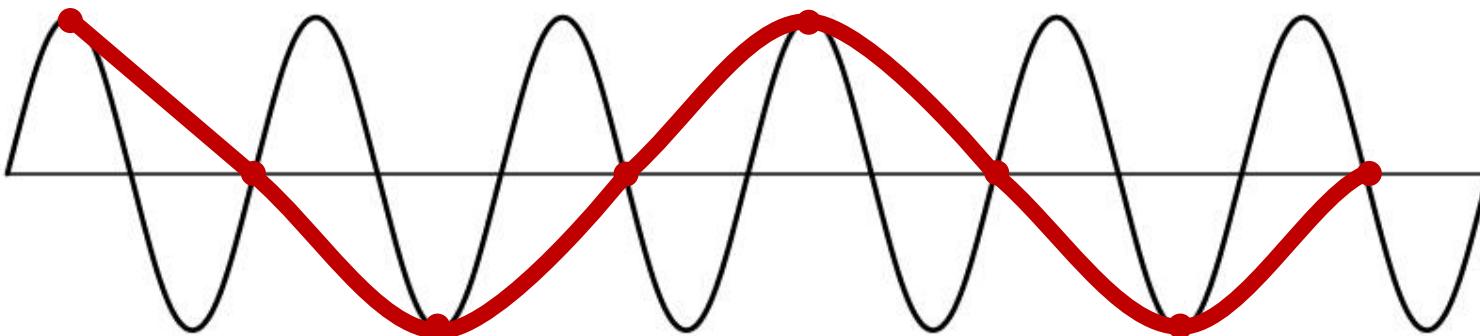
Nyquist Theorem

- Sampling rate should be equal to or greater than **twice the highest frequency** in the analog signal



Nyquist Theorem

- Sampling rate should be equal to or greater than **twice the highest frequency** in the analog signal



Antialiasing

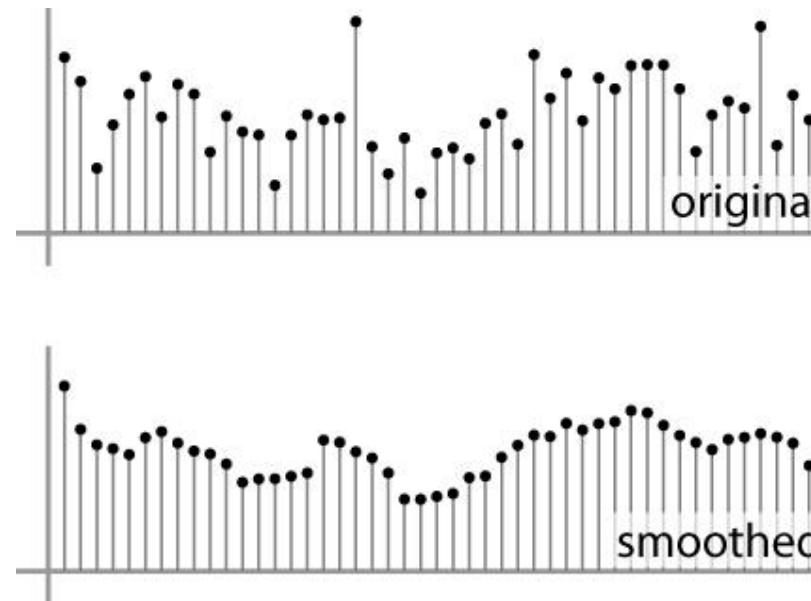
- **Fast-changing signals sampled too slowly**
- **What can we do about aliasing?**
 - **Sample more often**
 - **Join the Megapixel craze of the photo industry**
 - **Only shifts the problem to higher frequencies**
 - **Make the signal less “wiggly” (filtering)**
 - **Get rid of some high frequencies**
 - **Will lose information**
 - **But it's better than aliasing**

Outline

- **Sampling and aliasing**
- **Frequency domain**
- **Smoothing filters**
- **Antialiasing**
- **Other filters**

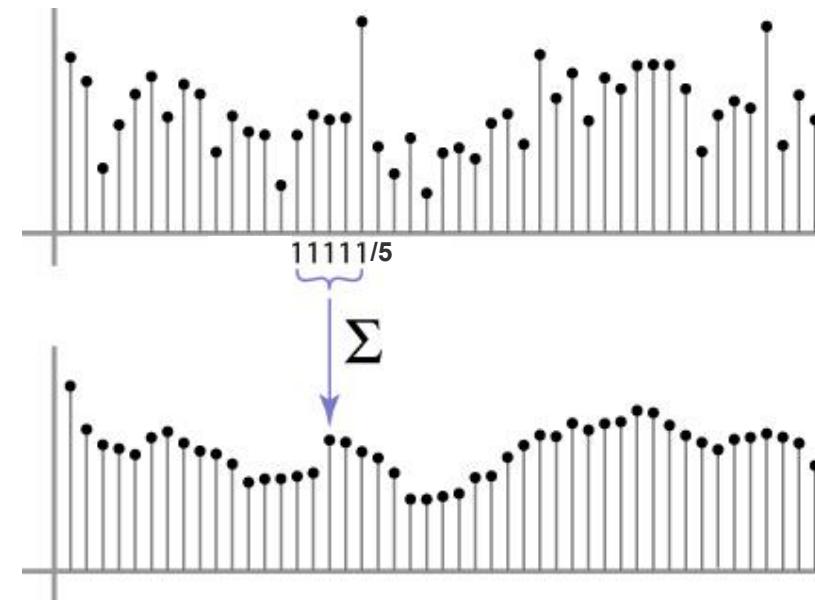
Moving Average

- Basic idea: define a new function by averaging over a sliding window



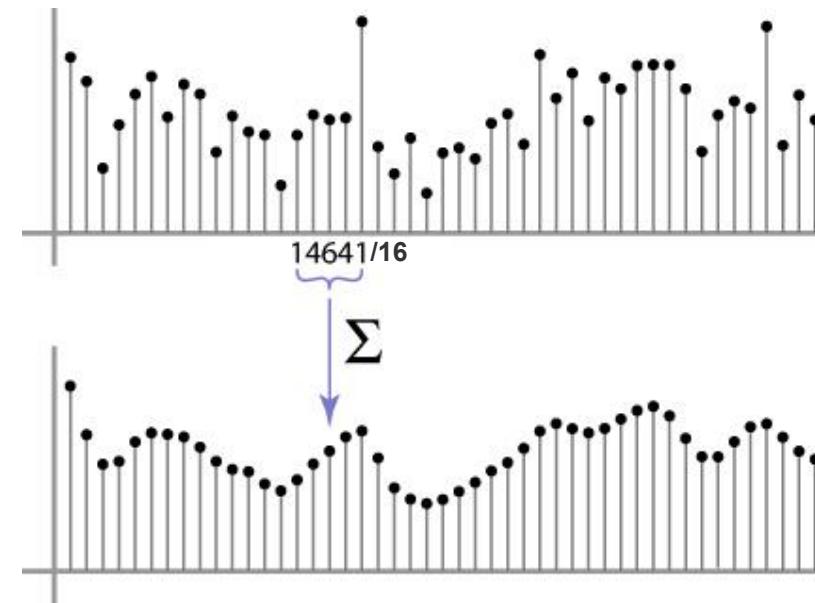
Weighted Moving Average

- Can add weights to our moving average
- Weights [1, 1, 1, 1, 1] / 5



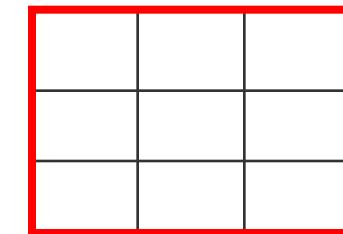
Weighted Moving Average

- Bell curve (gaussian-like) weights [1, 4, 6, 4, 1]/16



Weighted Moving Average In 2D

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

 $F[x, y]$  $H[u, v]$

Cross-correlation

- Let F be the image, H be the kernel (of size $2k+1 \times 2k+1$), and G be the output image

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v]F[i + u, j + v]$$

- H is called “filter”, “kernel”, or “mask”
- This is called a cross-correlation operation:

$$G = H \otimes F$$

Moving Average In 2D

□ What are the weights H?

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

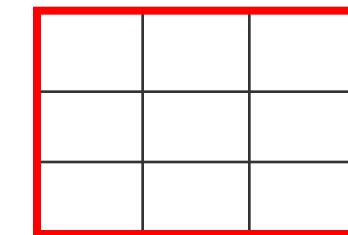
 $F[x, y]$  $H[u, v]$

Image filtering

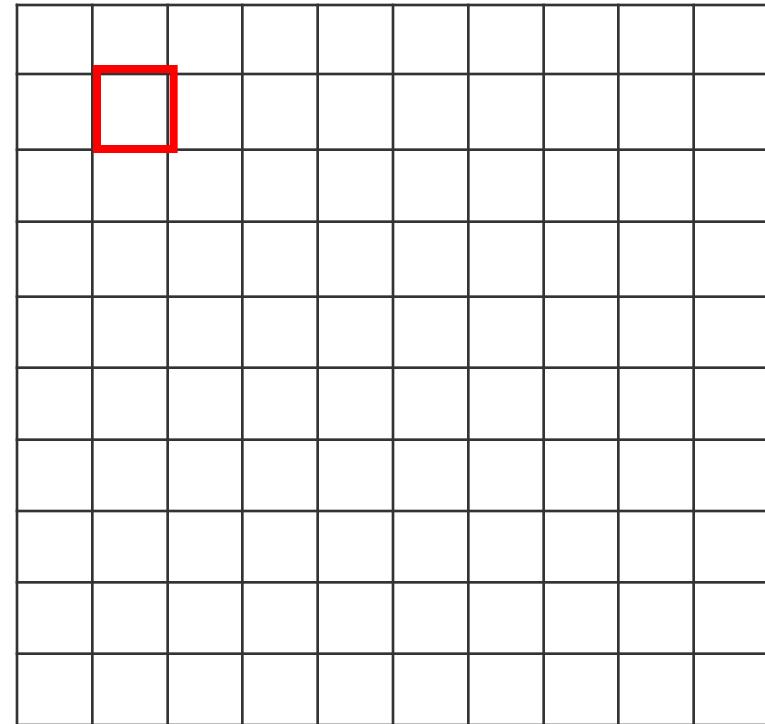
$$h[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$f[.,.]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$g[.,.]$



$$g[m, n] = \sum_{k,l} h[k, l] f[m + k, n + l]$$

Credit: S. Seitz

Image filtering

$$f[.,.]$$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$$g[.,.]$$

$$h[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

	0	10								

$$g[m, n] = \sum_{k,l} h[k, l] f[m + k, n + l]$$

Credit: S. Seitz

Image filtering

$f[.,.]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$g[.,.]$

			0	10	20					

$$h[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

Image filtering

$f[.,.]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$g[.,.]$

			0	10	20	30				

$$h[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

Image filtering

$$h[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$f[.,.]$

0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0

$g[.,.]$

	0	10	20	30	30						

Image filtering

$$h[\cdot, \cdot] \frac{1}{9}$$

$$f[.,.]$$

$$g[.,.]$$

Image filtering

 $f[.,.]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

 $g[.,.]$

	0	10	20	30	30					

$$h[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

Image filtering

$$h[\cdot, \cdot] \quad \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$$f[., .]$$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$$g[., .]$$

	0	10	20	30	30	30	20	10		
	0	20	40	60	60	60	40	20		
	0	30	60	90	90	90	60	30		
	0	30	50	80	80	90	60	30		
	0	30	50	80	80	90	60	30		
	0	20	30	50	50	60	40	20		
	10	20	30	30	30	30	20	10		
	10	10	10	0	0	0	0	0		

$$g[m, n] = \sum_{k,l} h[k, l] f[m + k, n + l]$$

Credit: S. Seitz

Box filter

□ What does it do?

- Replaces each pixel with an average of its neighborhood
- Achieve smoothing effect (remove sharp features)

$$h[\cdot, \cdot]$$

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

Cross-correlation vs. Convolution

- **Cross-correlation:**

$$G = H \otimes F$$

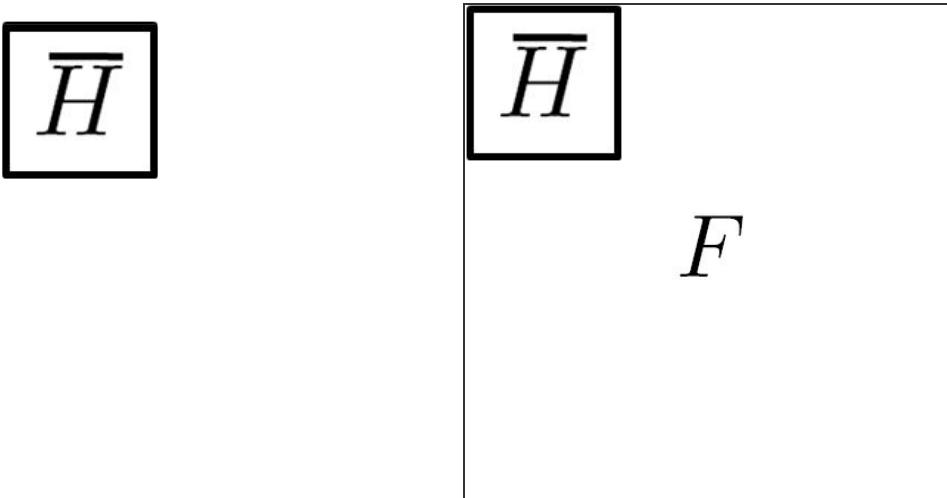
$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v]F[i + u, j + v]$$

- **Convolution = cross-correlation with a horizontally and vertically flipped filter**

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v]F[i - u, j - v]$$

$$G = H \star F$$

Convolution



Convolution is nice!

- Convolution is a multiplication-like operation
 - Commutative $a \star b = b \star a$
 - Associative $a \star (b \star c) = (a \star b) \star c$
 - Distributes over addition $a \star (b + c) = a \star b + a \star c$
 - Scalars factor out $\alpha a \star b = a \star \alpha b = \alpha(a \star b)$
 - Identity: unit impulse $e = [..., 0, 0, 1, 0, 0, ...]$ $a \star e = a$
- Conceptually no distinction between filter and signal

Convolution is nice!

- **Usefulness of associativity**
 - often apply several filters one after another
 - $((((a * b_1) * b_2) * b_3) = a * (b_1 * b_2 * b_3)$

Gaussian filtering

- A Gaussian kernel gives less weight to pixels further from the center of the window

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$F[x, y]$

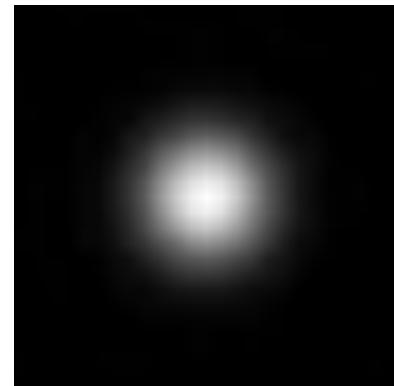
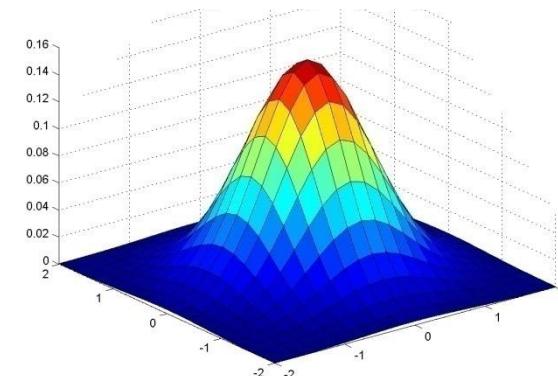
$$\frac{1}{16} \begin{matrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{matrix}$$

$H[u, v]$

Gaussian filter

- This kernel is an approximation of a Gaussian function:

$$G_\sigma = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

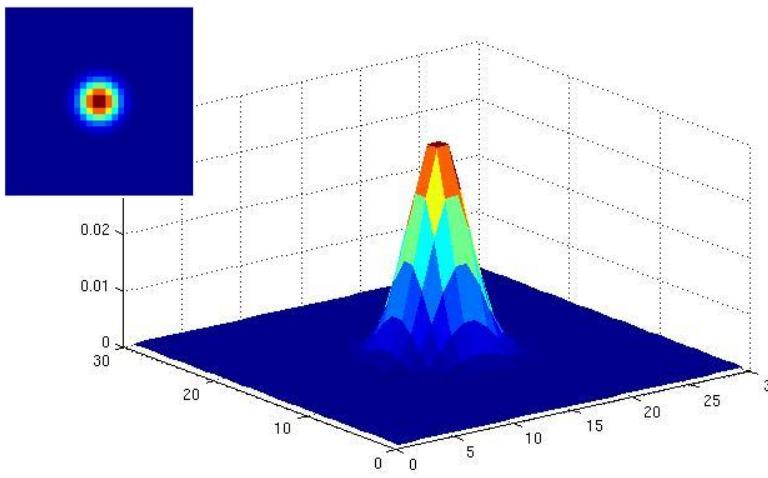


0.003	0.013	0.022	0.013	0.003
0.013	0.059	0.097	0.059	0.013
0.022	0.097	0.159	0.097	0.022
0.013	0.059	0.097	0.059	0.013
0.003	0.013	0.022	0.013	0.003

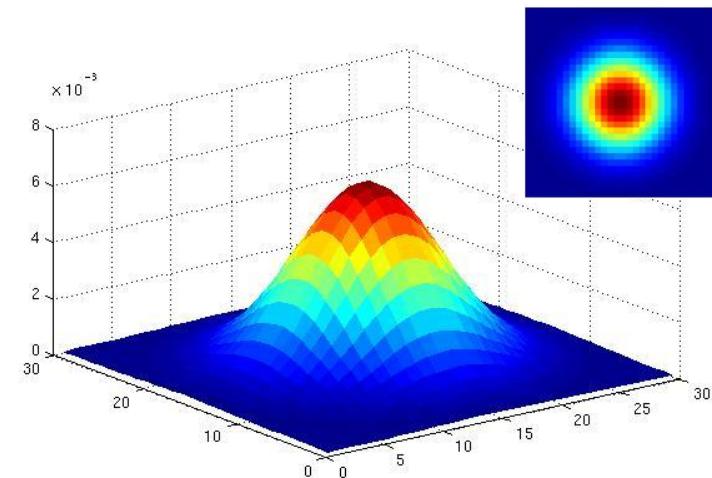
$5 \times 5, \sigma = 1$

Gaussian Kernel

- Standard deviation σ : determines extent of smoothing



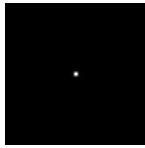
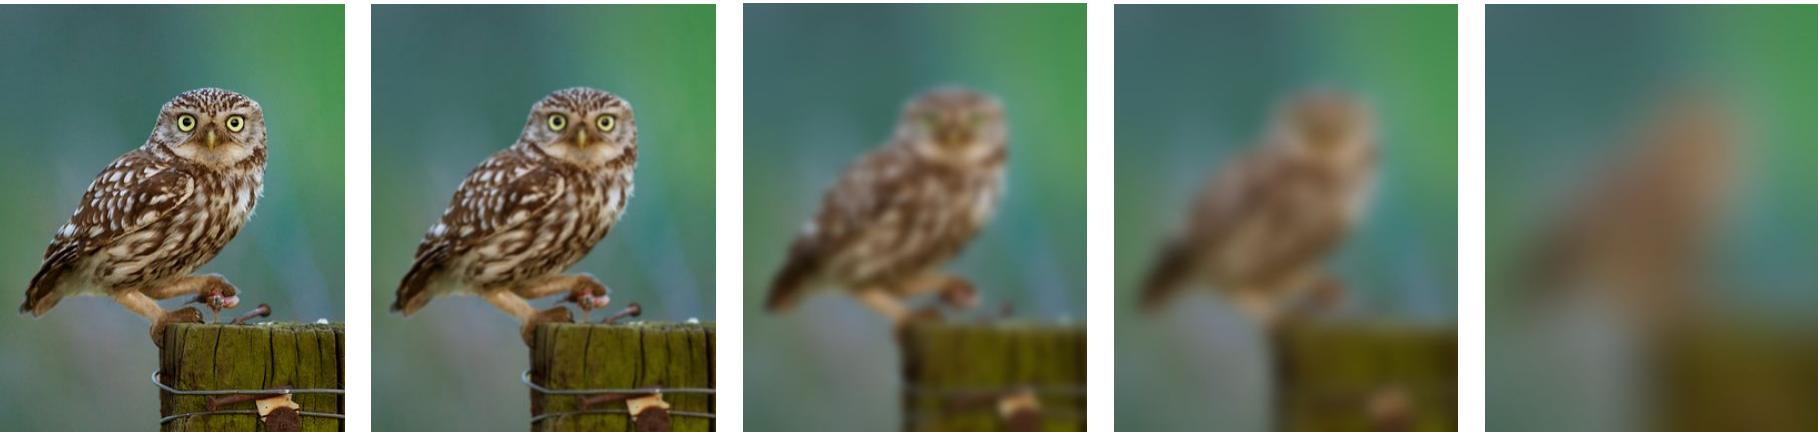
$\sigma = 2$ with 30×30 kernel



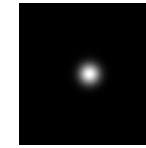
$\sigma = 5$ with 30×30 kernel

$$G_\sigma = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

Gaussian filters



$\sigma = 1$ pixel



$\sigma = 5$ pixels



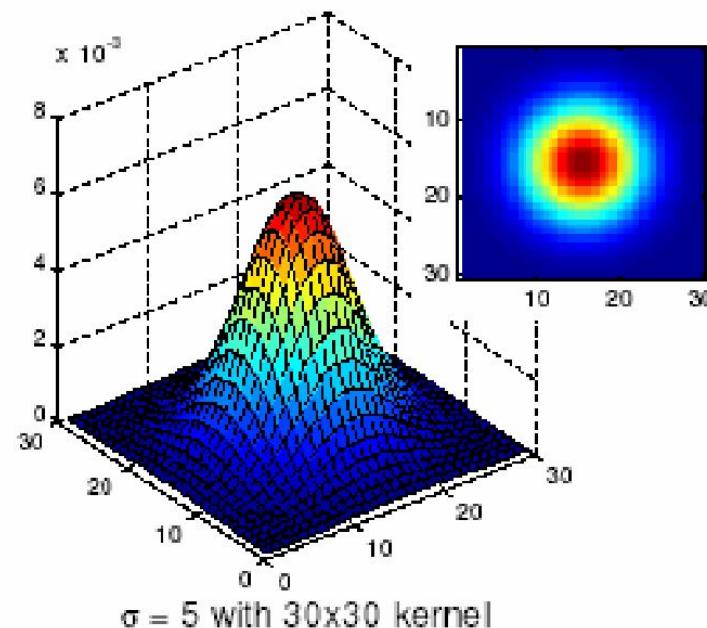
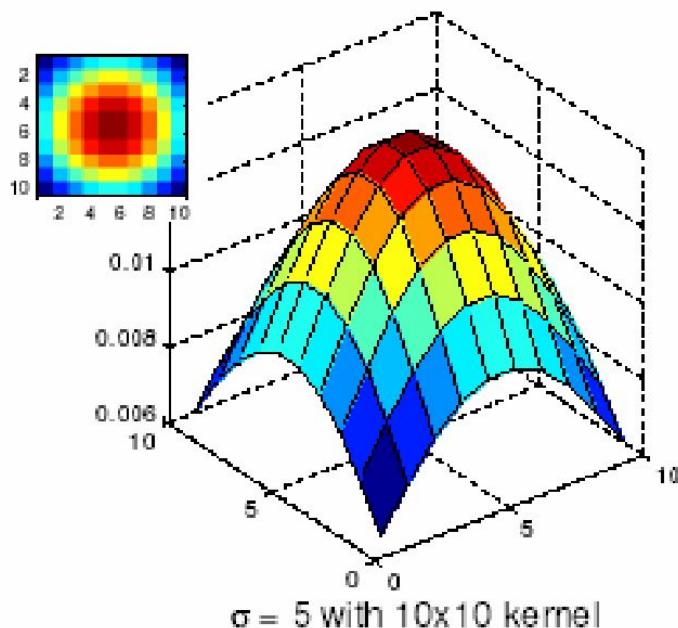
$\sigma = 10$ pixels



$\sigma = 30$ pixels

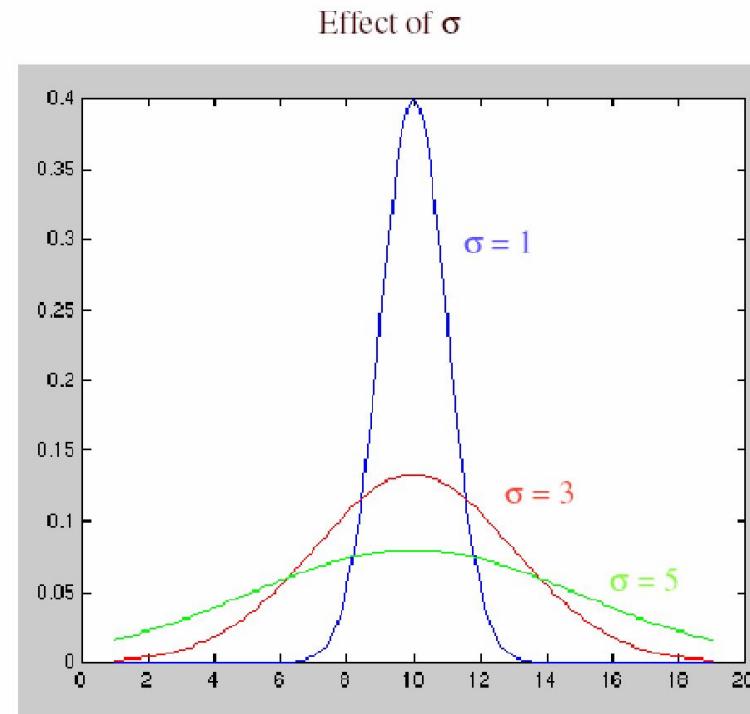
Choosing kernel width

- The Gaussian function has infinite support, but discrete filters use finite kernels



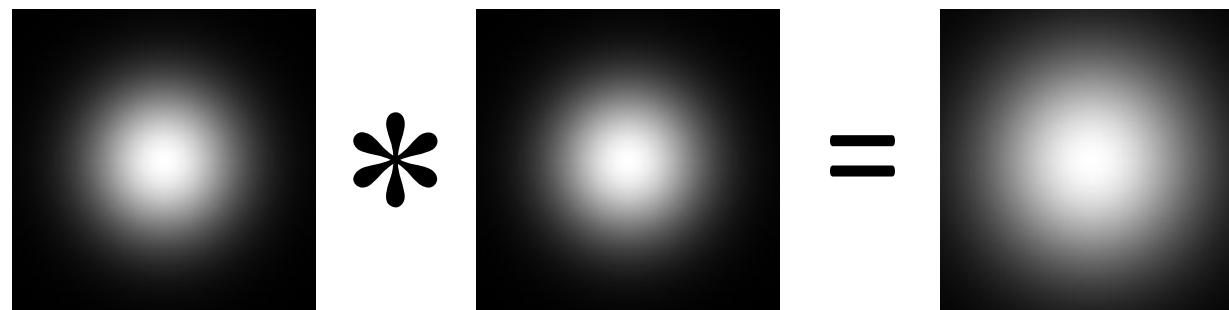
Practical matters

- How big should the filter be?
- Values at edges should be near zero
- Rule of thumb for Gaussian: set filter half-width to about 3σ

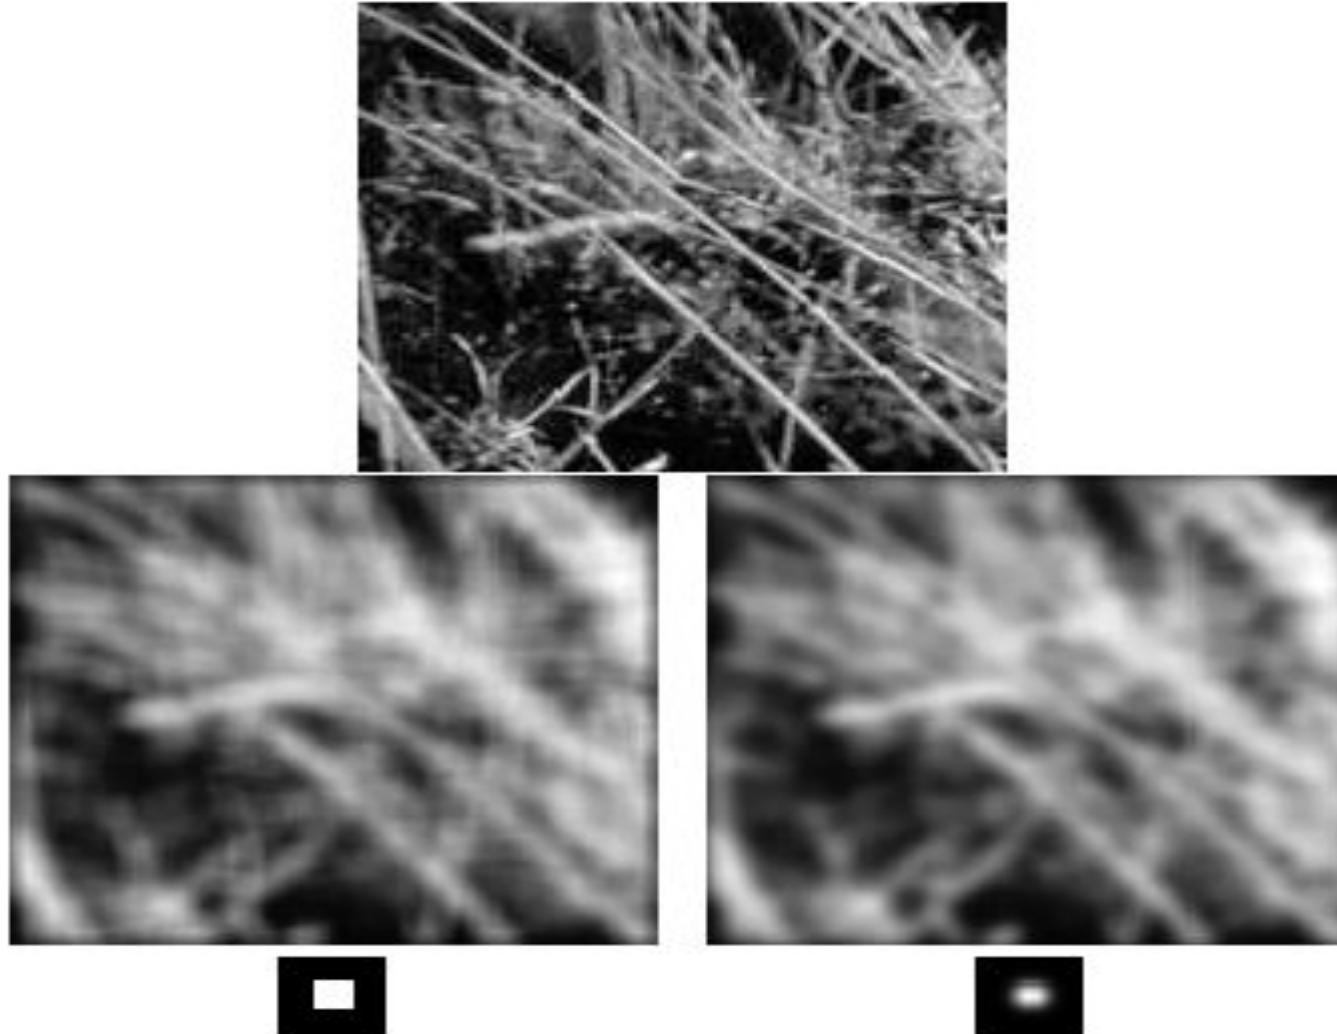


Gaussian and convolution

- Removes “high-frequency” components from the image (low-pass filter)
- Convolution with self is another Gaussian
 - Convolving twice with Gaussian kernel of std. σ is equal to convolving once with kernel of std. $\sigma\sqrt{2}$



Box vs. Gaussian filtering

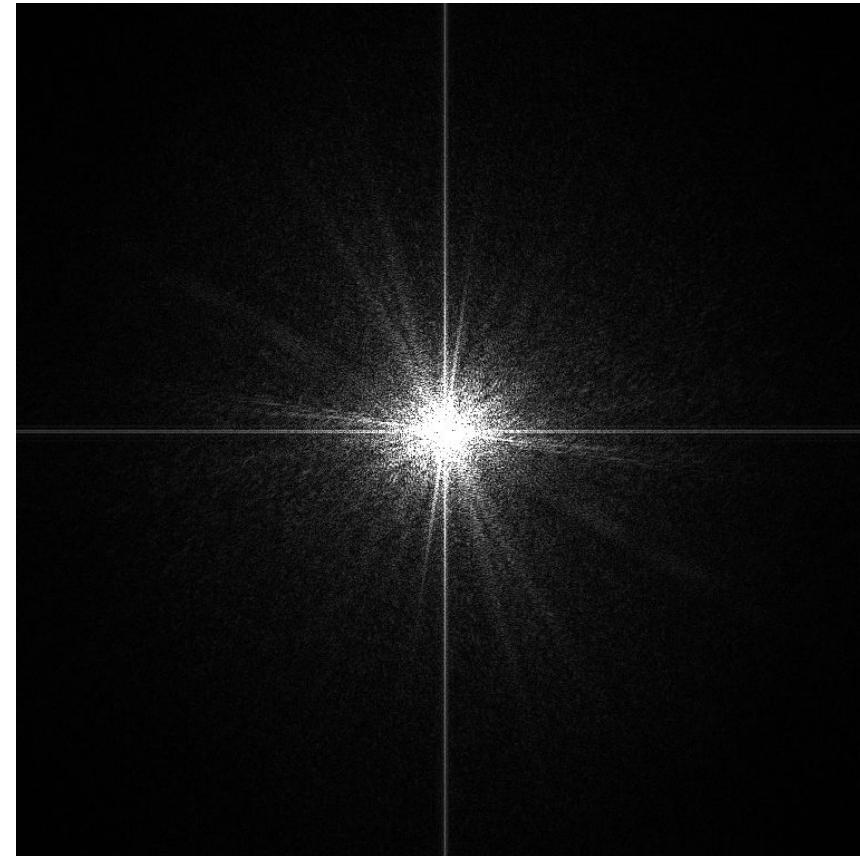


Filtering in Frequency Domain

Visualizing Image Frequency Content

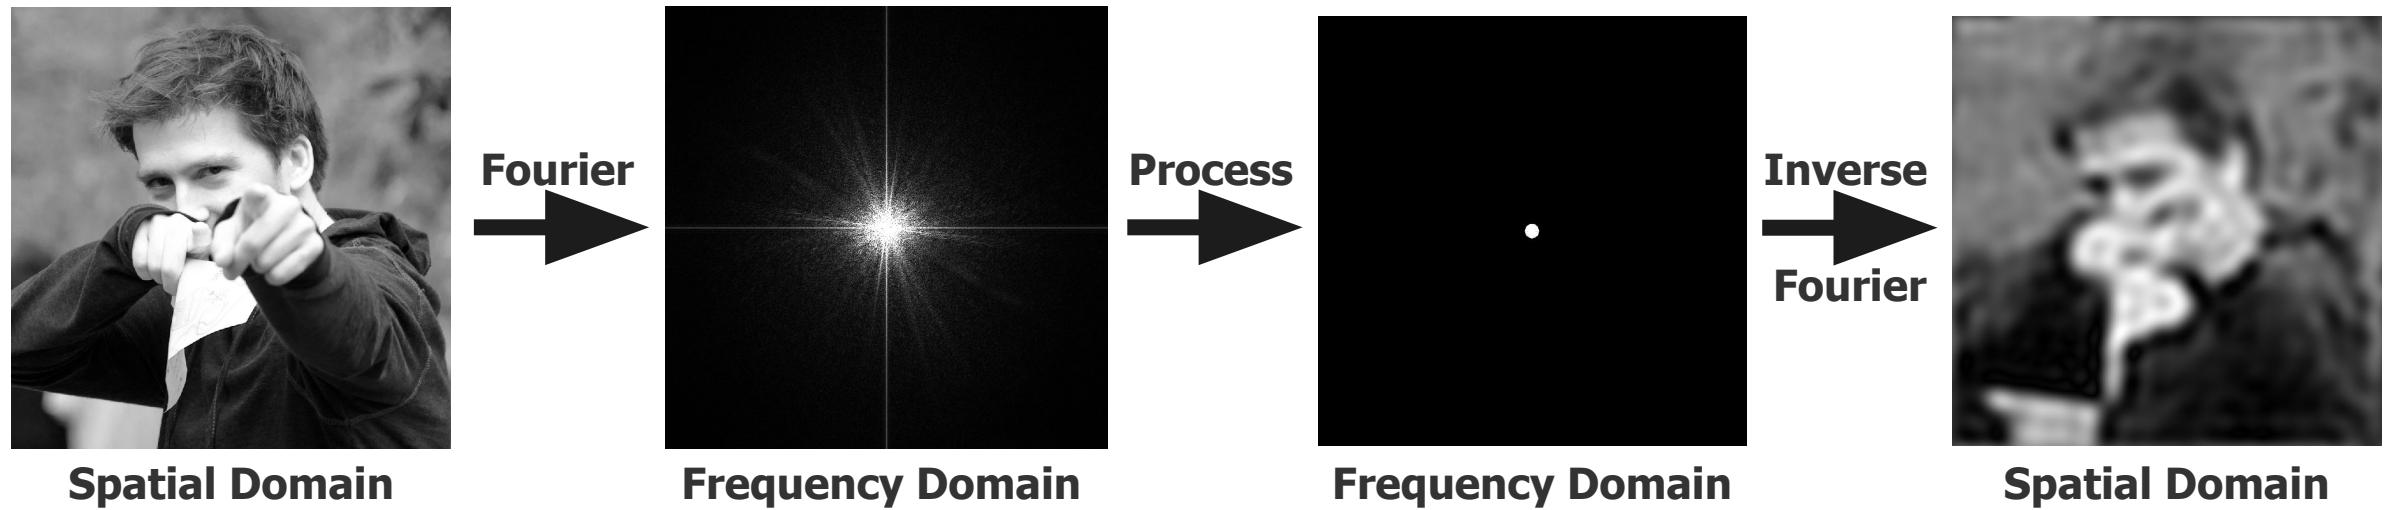


Spatial Domain



Frequency Domain

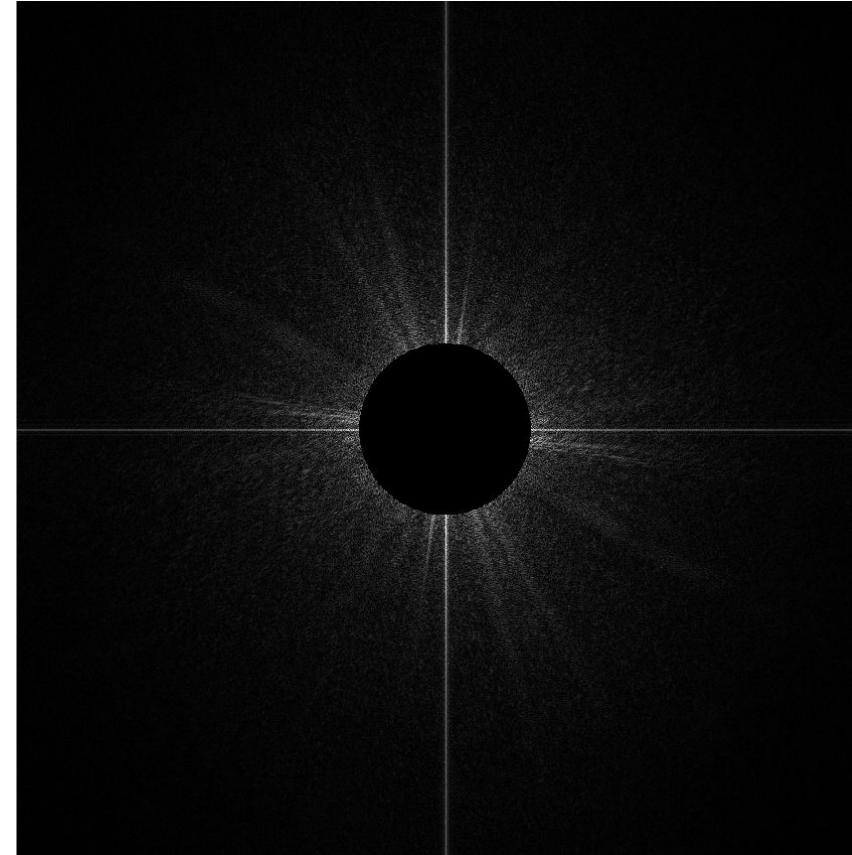
Filtering in the Frequency Domain



Filter Out Low Frequencies Only (Edges)



Spatial Domain

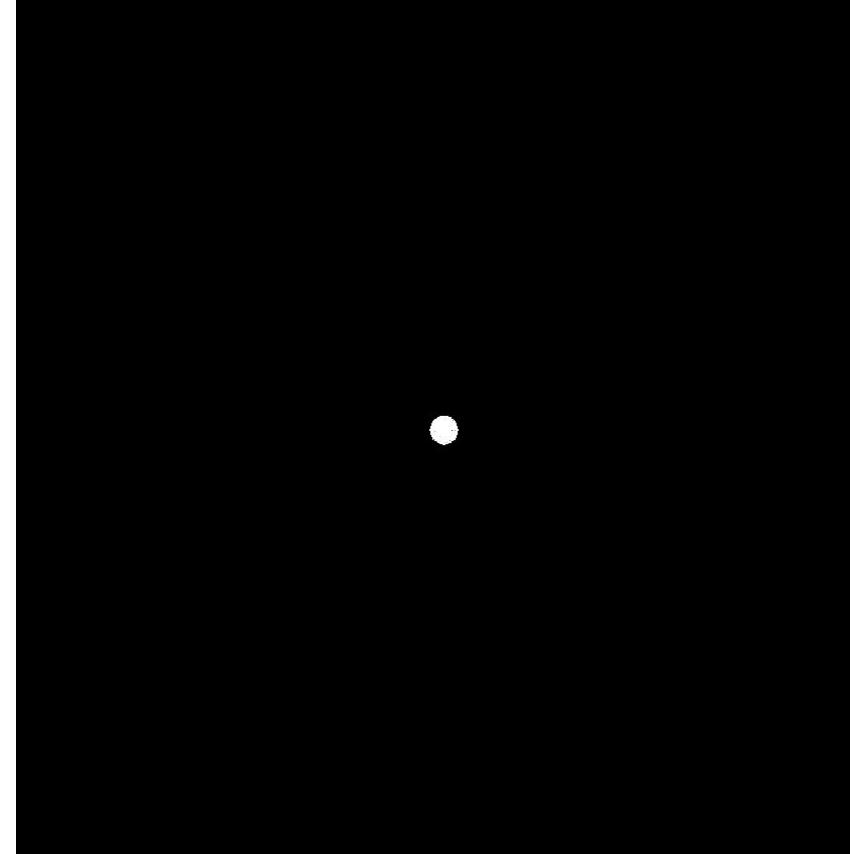


Frequency Domain

Filter Out High Frequencies (Blur)



Spatial Domain

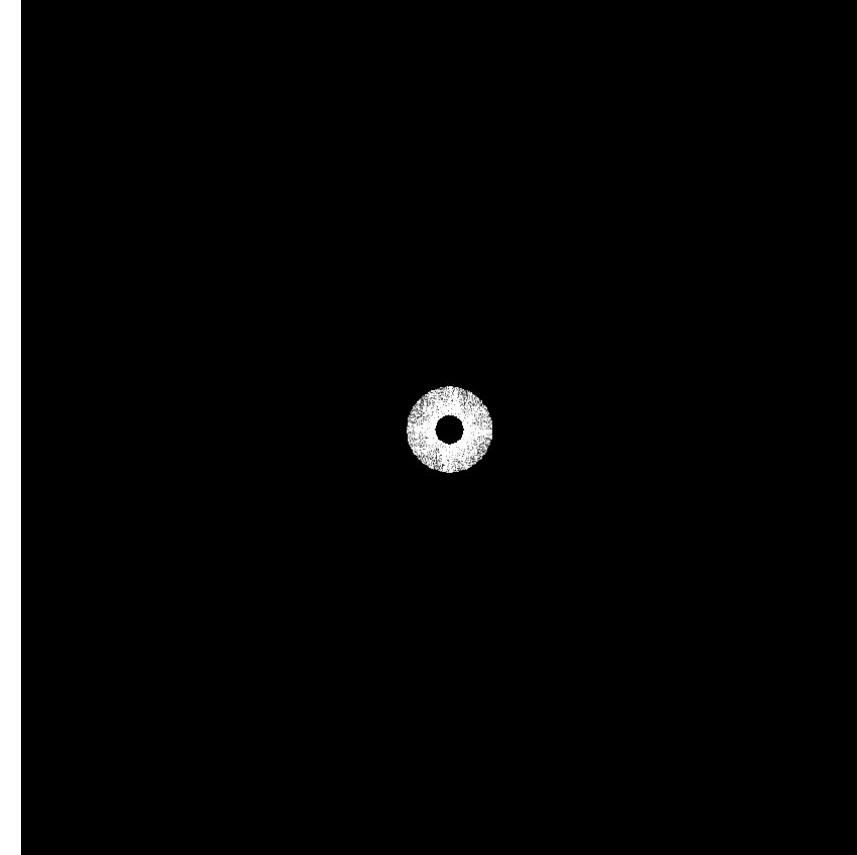


Frequency Domain

Filter Out Low and High Frequencies



Spatial Domain

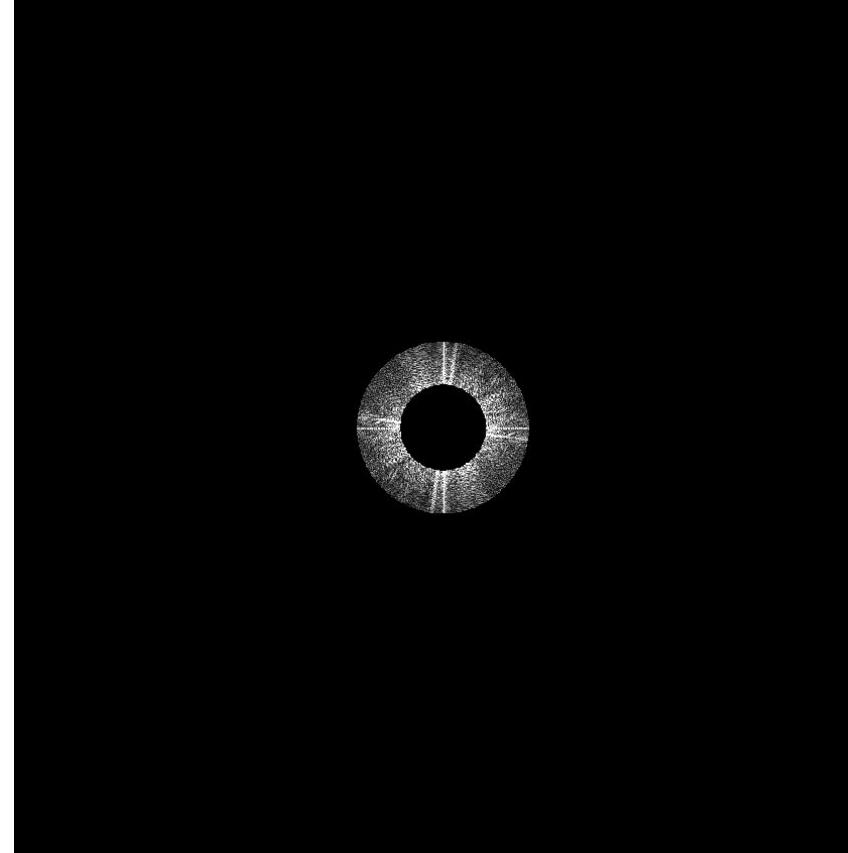


Frequency Domain

Filter Out Low and High Frequencies



Spatial Domain



Frequency Domain

Relationship between the two

The Convolution Theorem

- The Fourier transform of the convolution of two functions is the product of their Fourier transforms

$$\mathcal{F}[g * h] = \mathcal{F}[g]\mathcal{F}[h]$$

- Convolution in spatial domain is equivalent to multiplication in frequency domain!

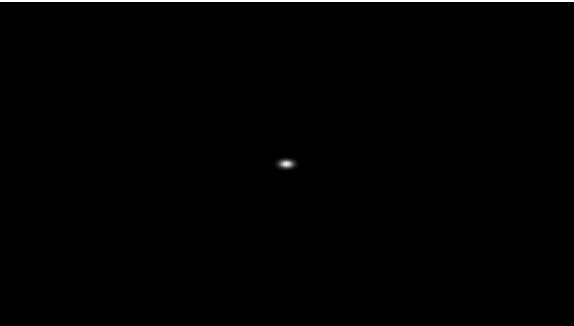
2D convolution theorem example

$f(x,y)$



*

$h(x,y)$



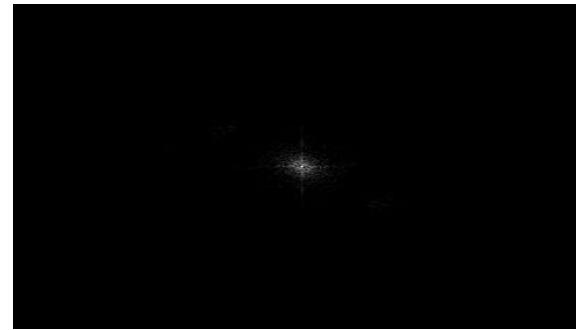
↓

$g(x,y)$



×

$|F(s_x, s_y)|$

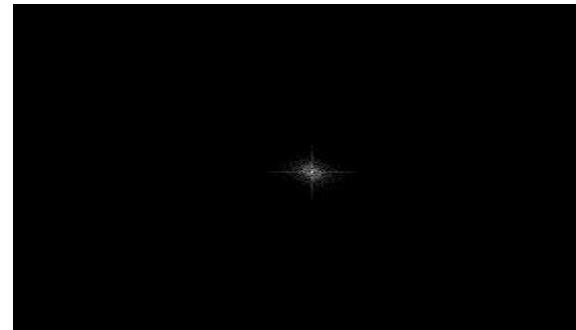


↓

$|H(s_x, s_y)|$



$|G(s_x, s_y)|$

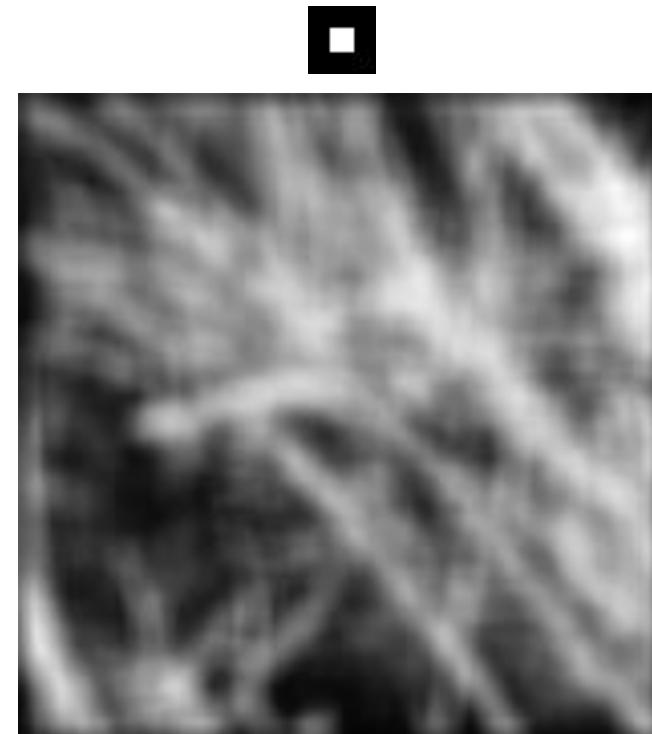


Box vs. Gaussian filtering

- Why does the Gaussian give a nice smooth image, but the square filter give edgy artifacts?

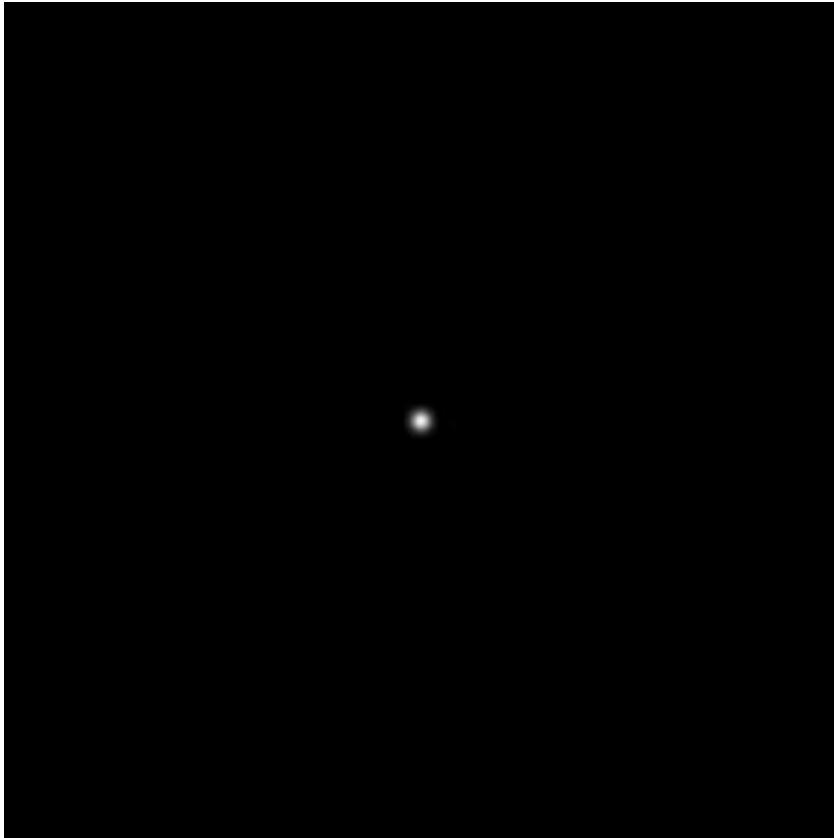


Gaussian

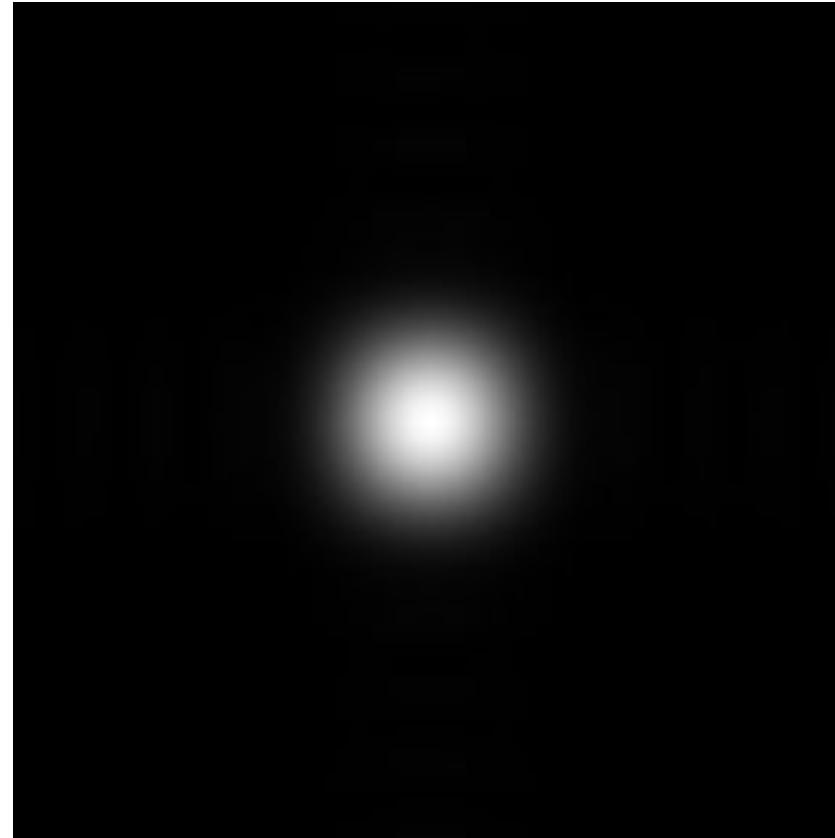


Box filter

Gaussian filter

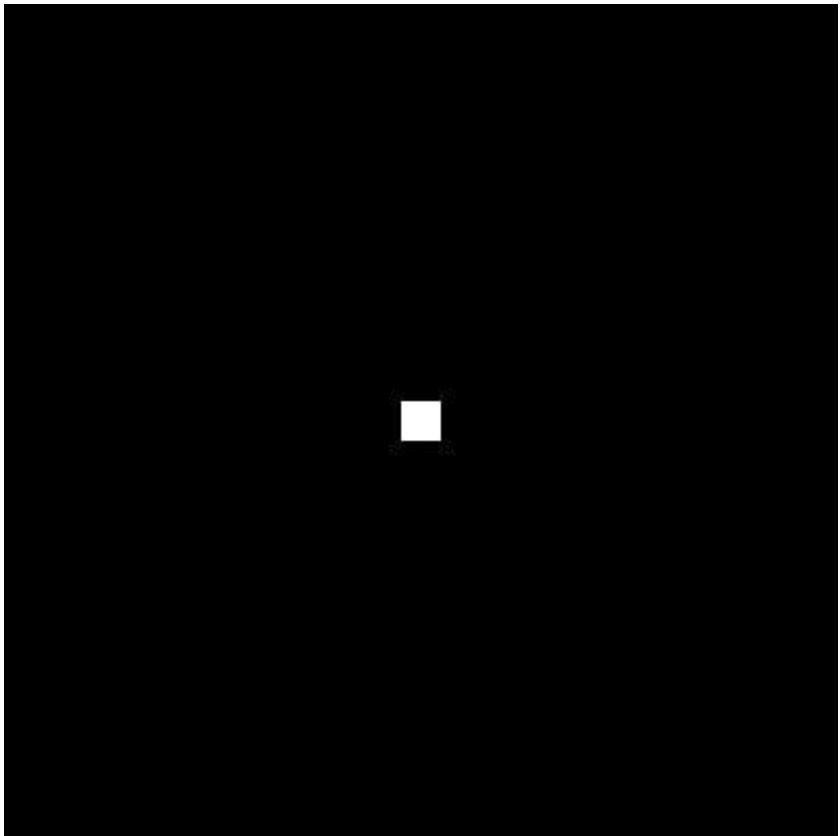


Spatial
Domain

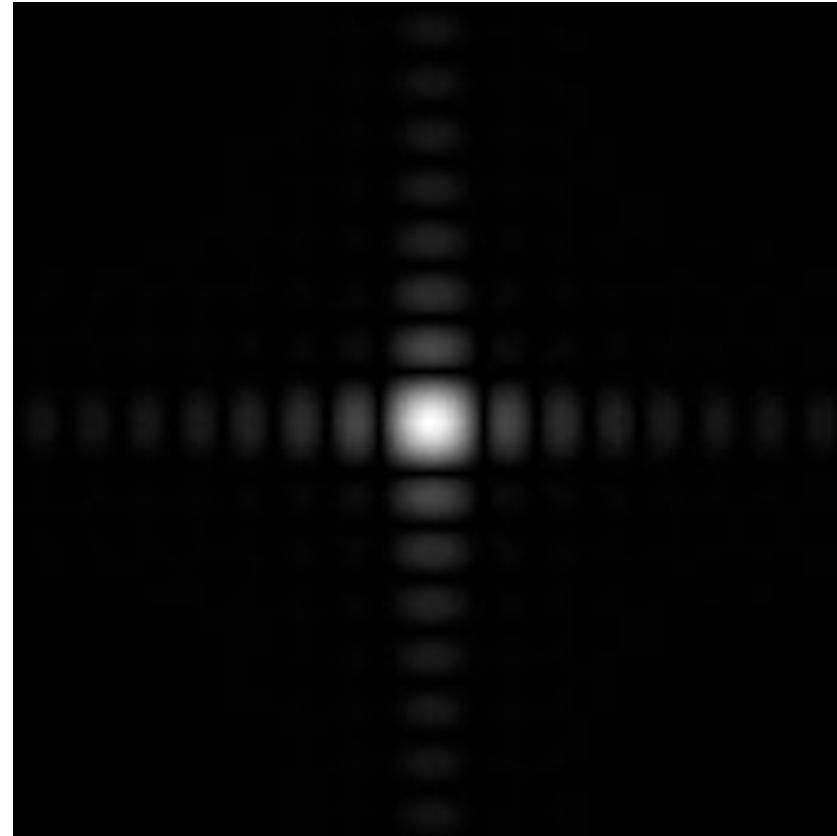


Frequency
Domain

Box filter



Spatial
Domain

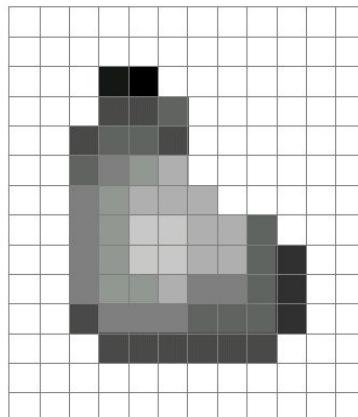
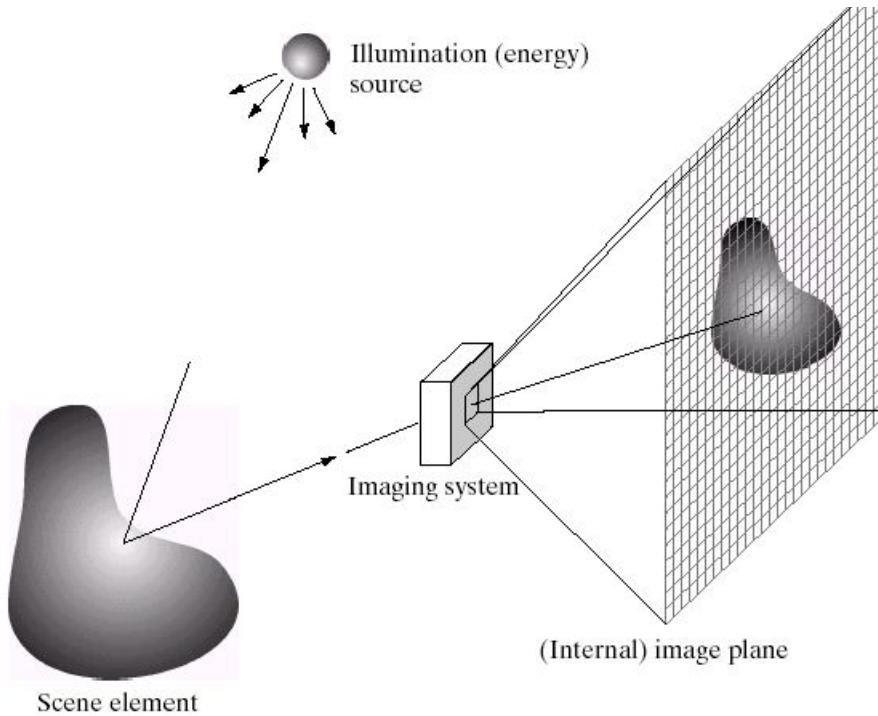


Frequency
Domain

Outline

- **Sampling and aliasing**
- **Frequency domain**
- **Smoothing filters**
- **Antialiasing**
- **Other filters**

Image formation involves sampling



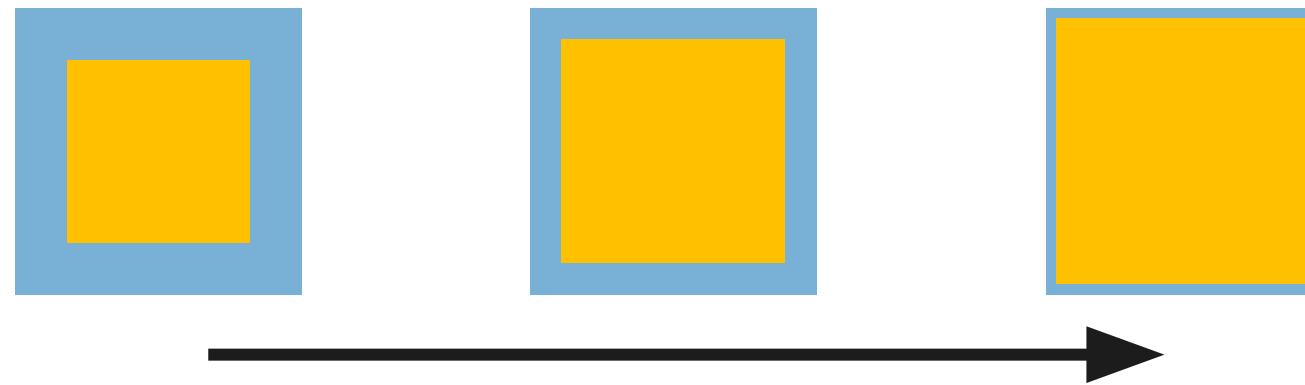
Antialiasing

- **Option 1**
 - **Filter the scene**
 - **Sample the filtered scene at each pixel**

- **Option 2**
 - **Average the scene under each pixel**

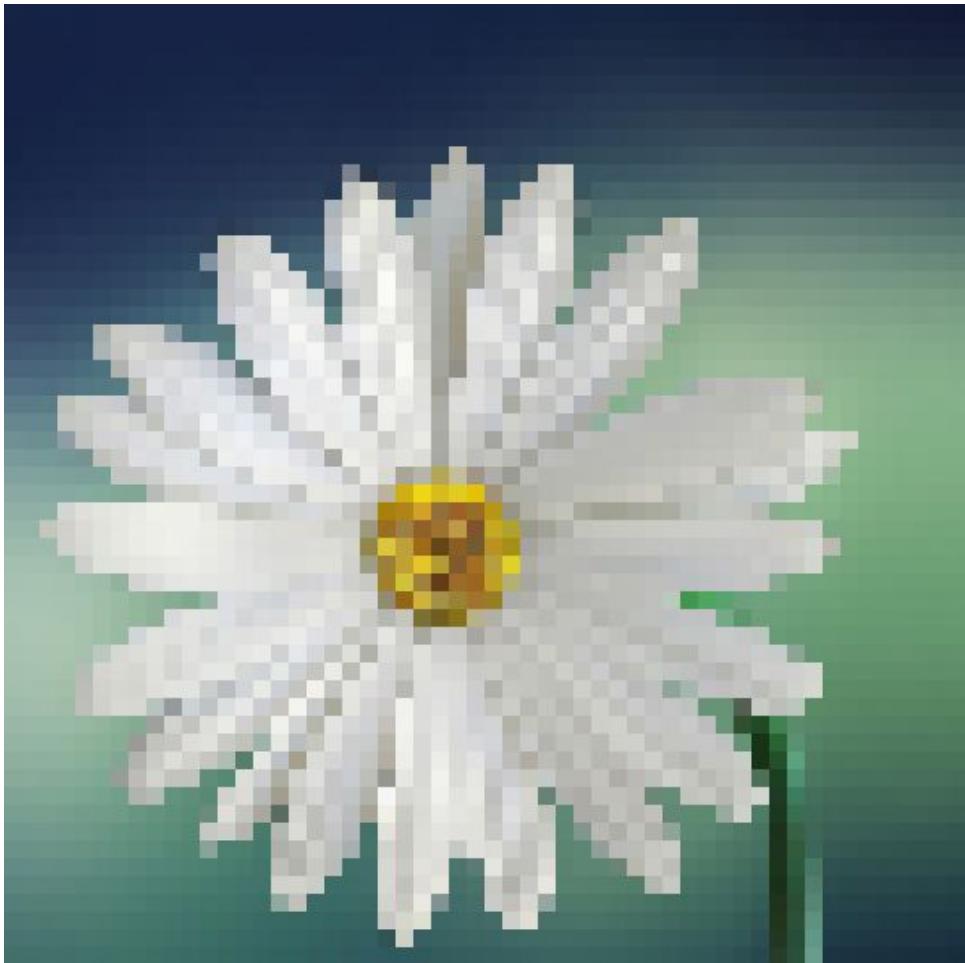
Digital sensors

- **Fill factor: ratio of light sensitive area to total pixel area**

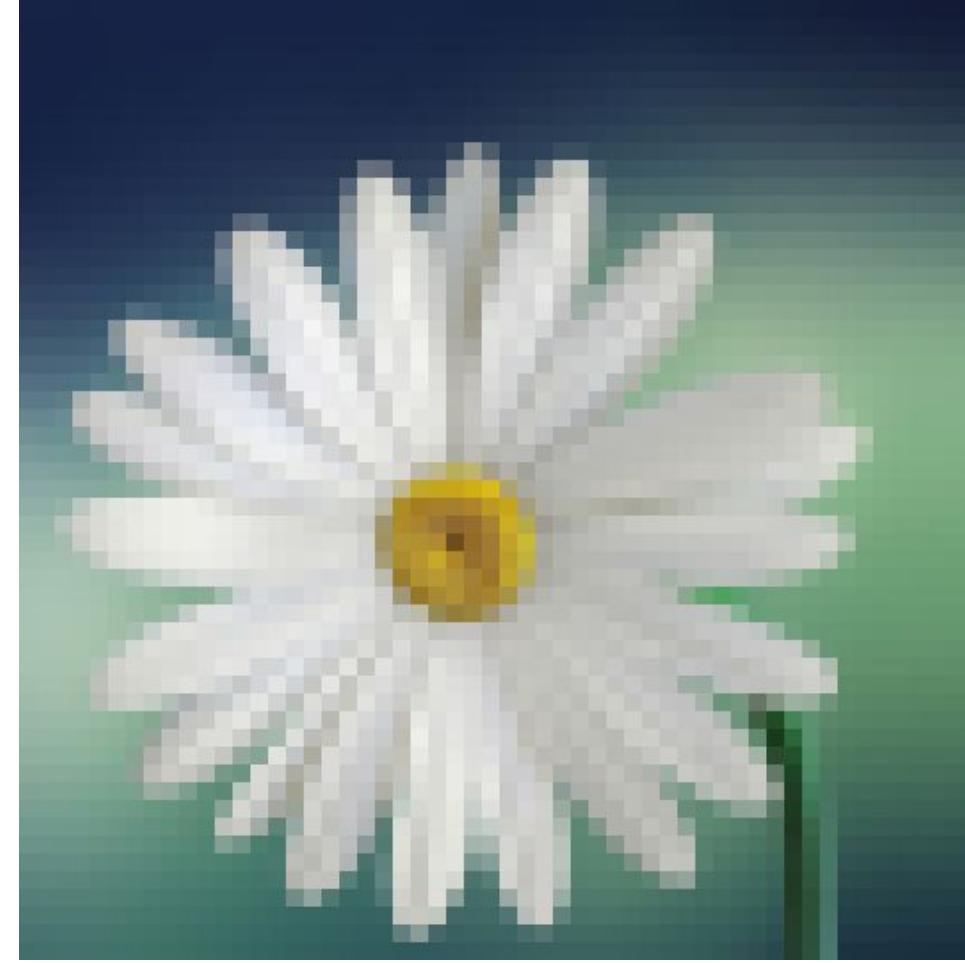


- Pixel area
- Light sensitive area

Fill factor effect



4%



100%

Outline

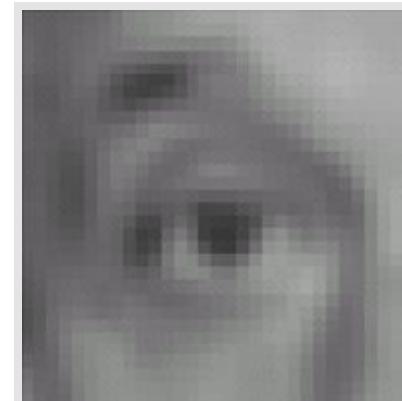
- **Sampling and aliasing**
- **Frequency domain**
- **Smoothing filters**
- **Antialiasing**
- **Other filters**

Linear filters: examples



Original

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} =$$



Blur
(with a mean filter)

Practice with linear filters

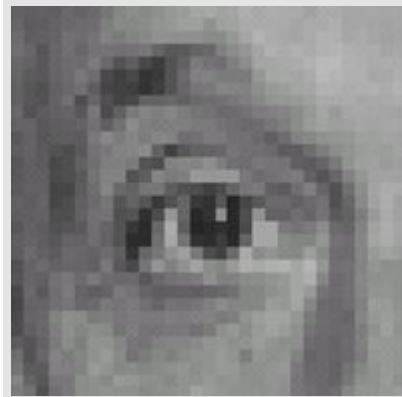


Original

0	0	0
0	1	0
0	0	0

?

Practice with linear filters



Original

0	0	0
0	1	0
0	0	0



Filtered
(no change)

Practice with linear filters



Original

0	0	0
0	0	1
0	0	0

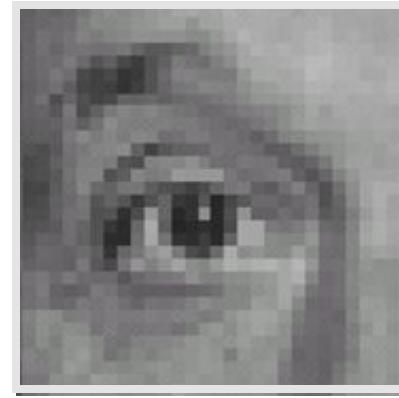
?

Practice with linear filters



Original

0	0	0
0	0	1
0	0	0



Shifted left
By 1 pixel

Linear filters: examples



Original

$$\left(\begin{array}{ccc} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{array} \right) - \frac{1}{9} \left(\begin{array}{ccc} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{array} \right) = \text{Sharpening filter output}$$
The result of applying the sharpening filter to the original image. The edges are more pronounced and the overall contrast is higher, making the features stand out more.

Sharpening filter
(accentuates edges)

Filtering – sharpening

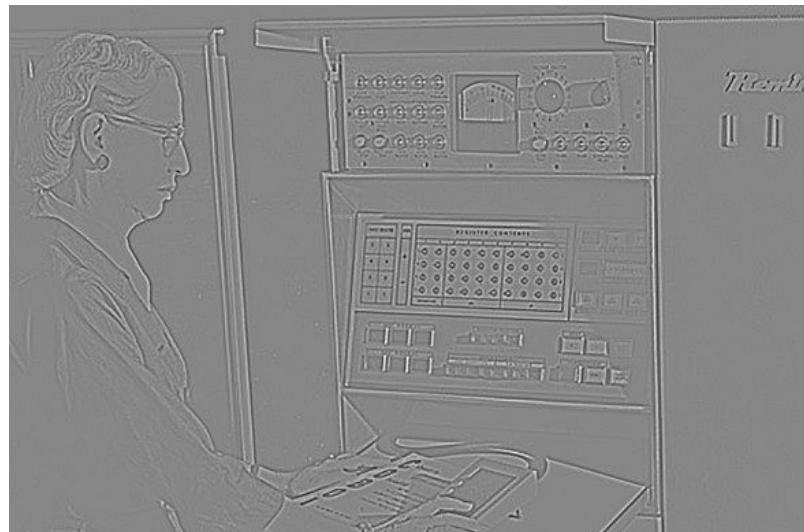
Image



Smoothed



Details

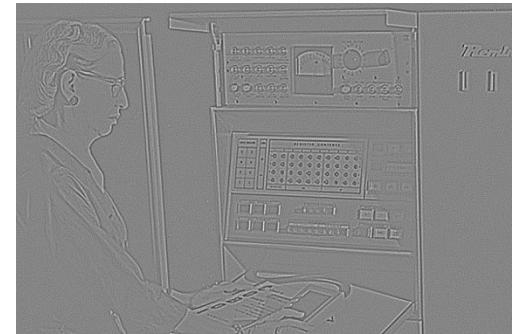


Filtering – Sharpening

Image



Details



$+ \alpha$

“Sharpened” $\alpha=1$

=

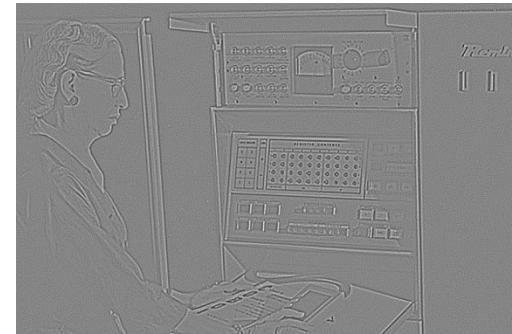


Filtering – Sharpening

Image



Details



$+ \alpha$

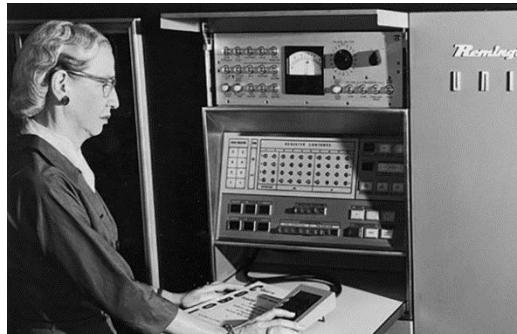
“Sharpened” $\alpha=0$

=

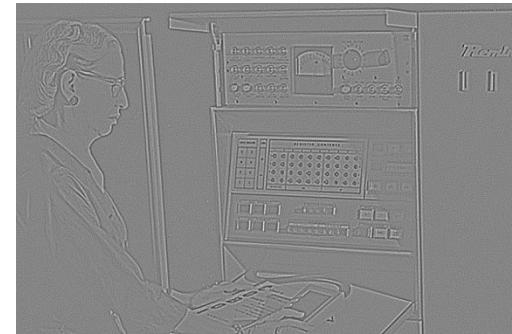


Filtering – Sharpening

Image



Details



$+ \alpha$

“Sharpened” $\alpha=2$

=

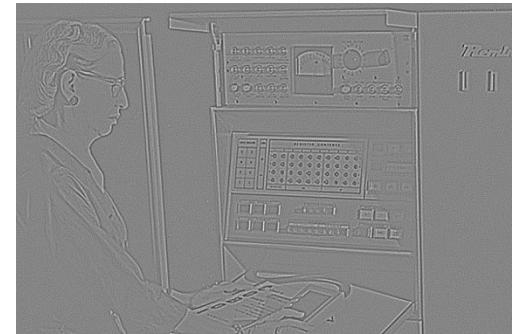


Filtering – Sharpening

Image



Details



$+ \alpha$

“Sharpened” $\alpha=0$

=

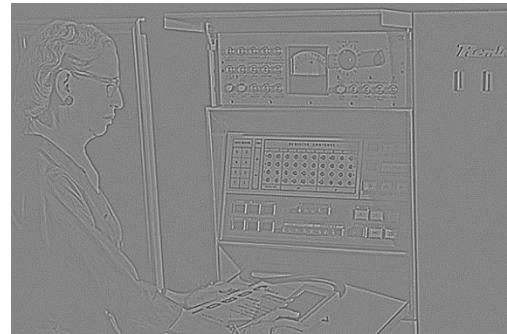


Filtering – Extreme Sharpening

Image



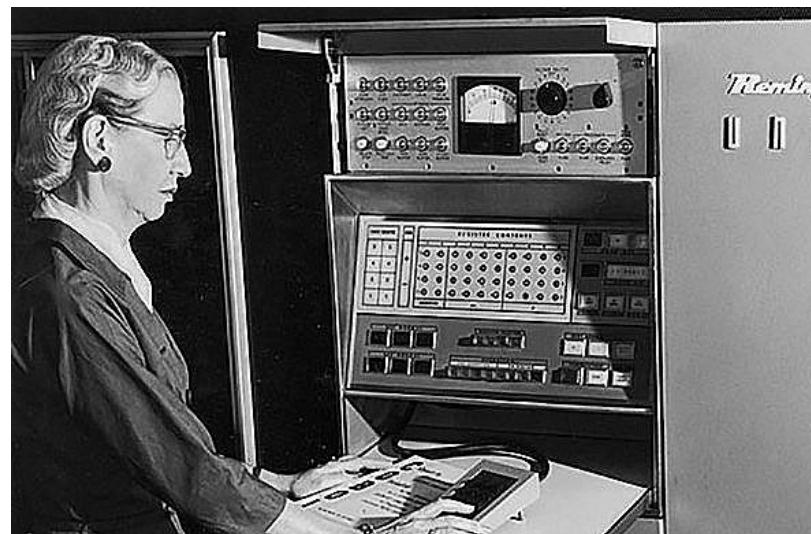
Details



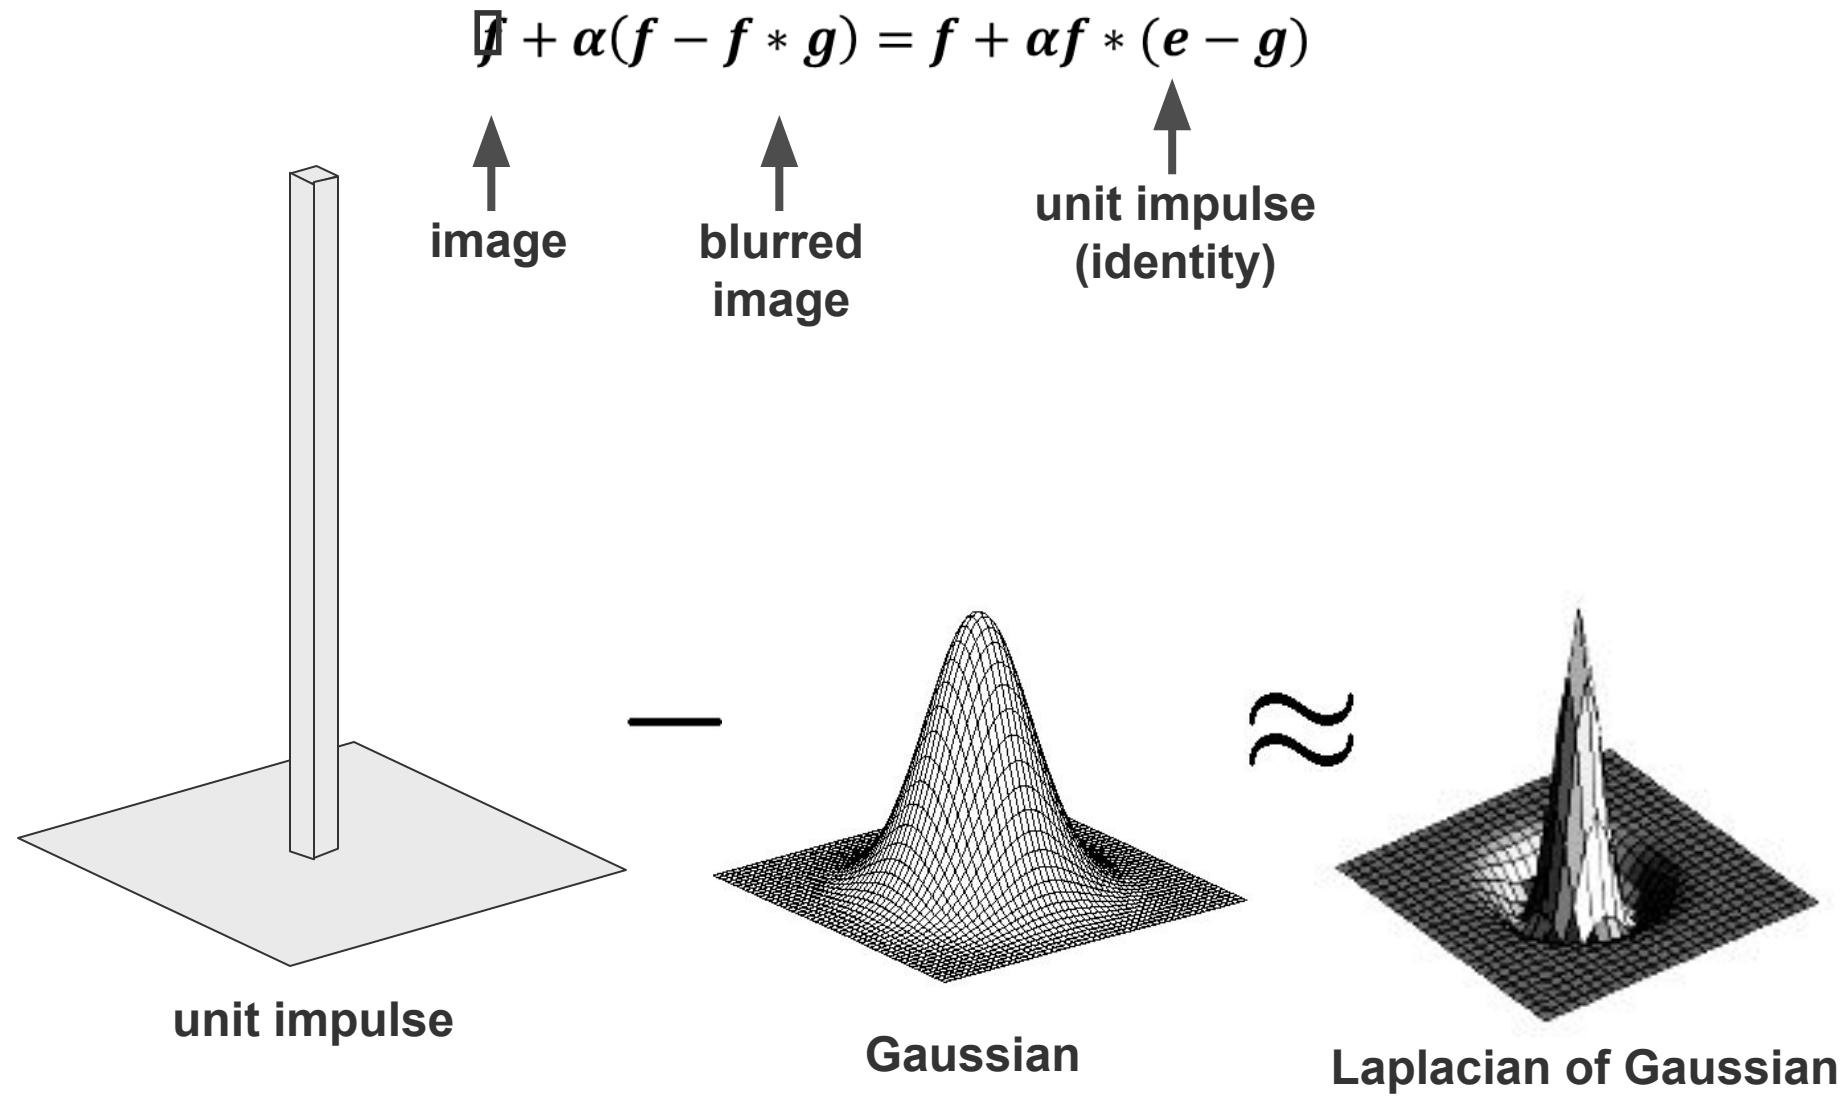
$+ \alpha$

“Sharpened” $\alpha=10$

=



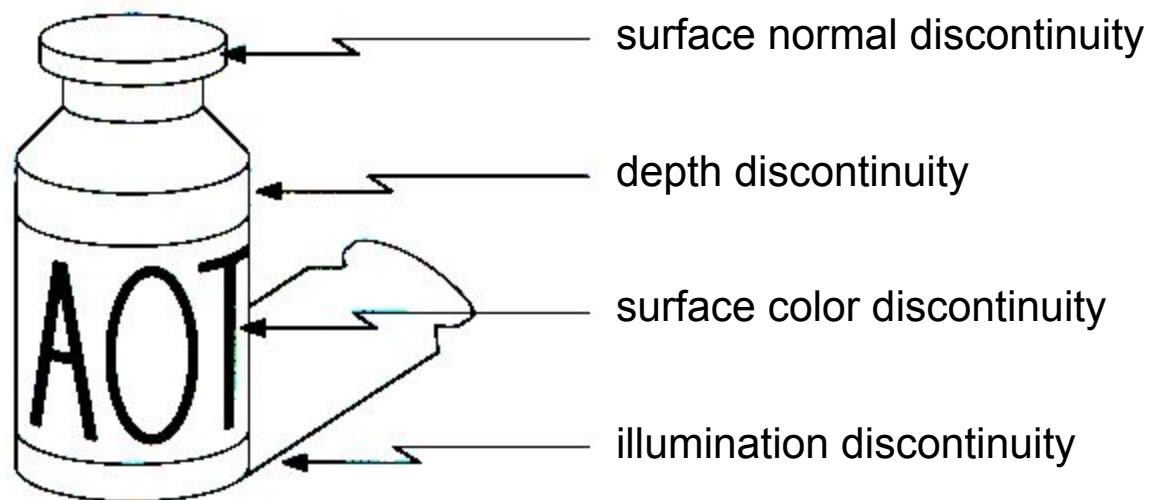
Unsharp mask filter



Edges



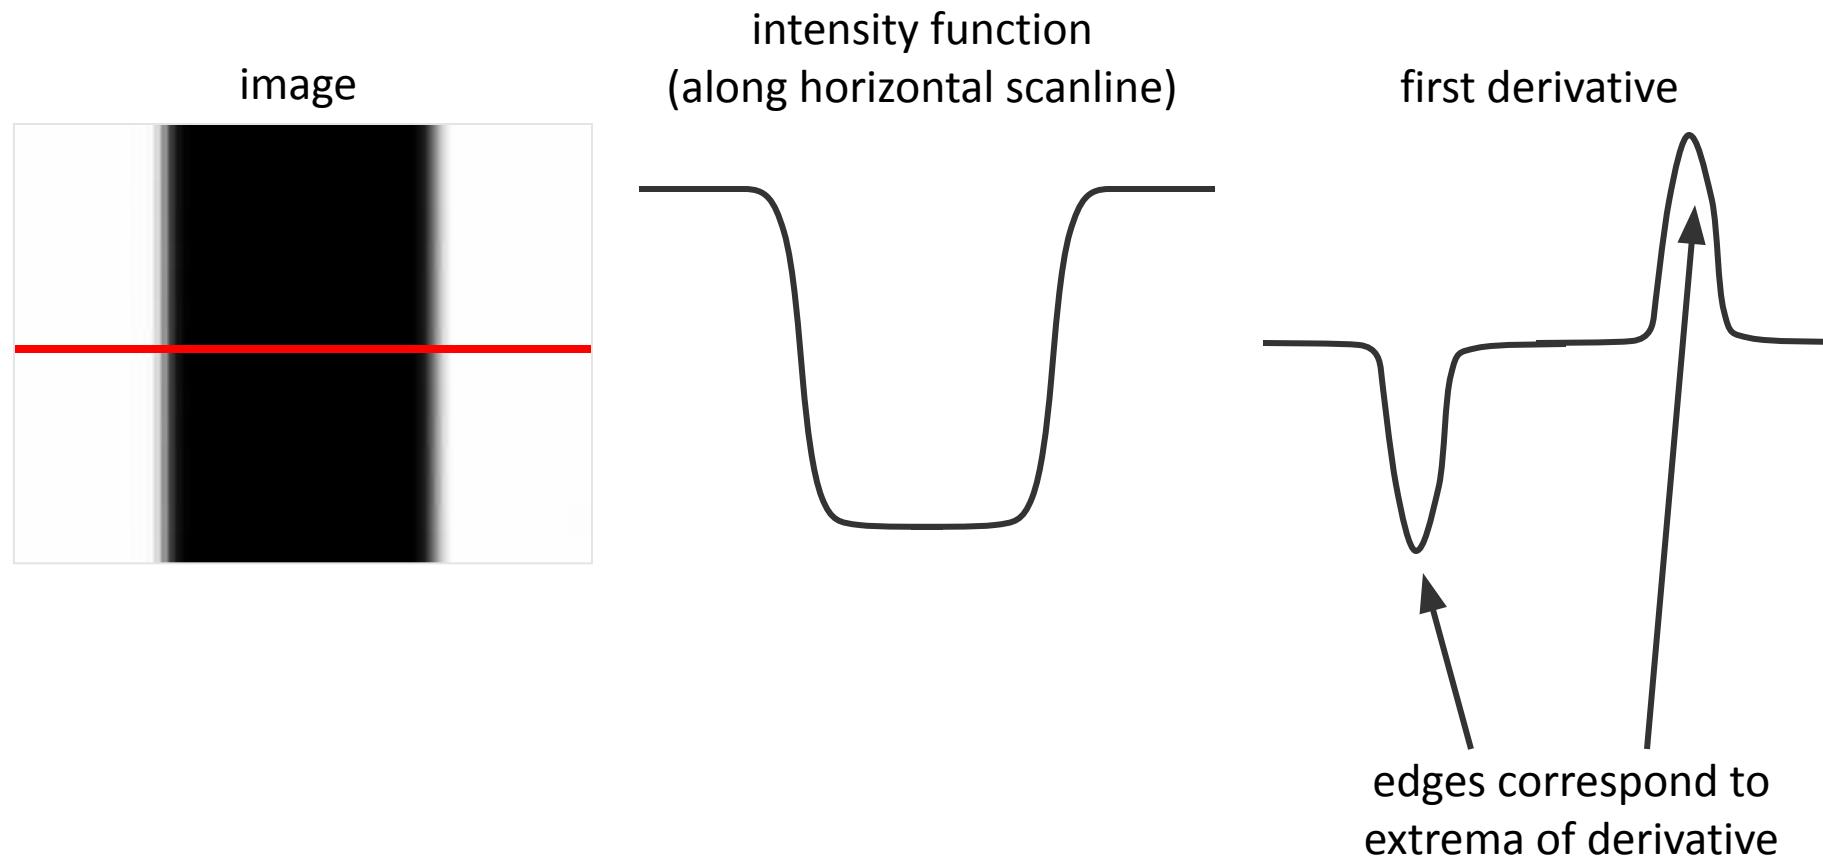
Origin of edges



Edges are caused by a variety of factors

Characterizing edges

- An edge is a place of rapid change in the image intensity function



Partial derivatives with convolution

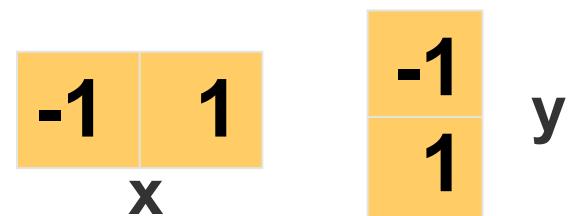
- For 2D function $f(x,y)$, the partial derivative is:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon}$$

- For discrete data, we can approximate using finite differences:

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x + 1, y) - f(x, y)}{1}$$

- To implement above as convolution, what would be the associated filter?

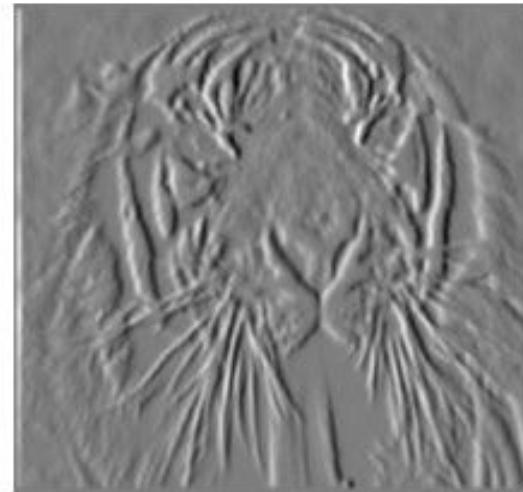


Partial derivatives of an image



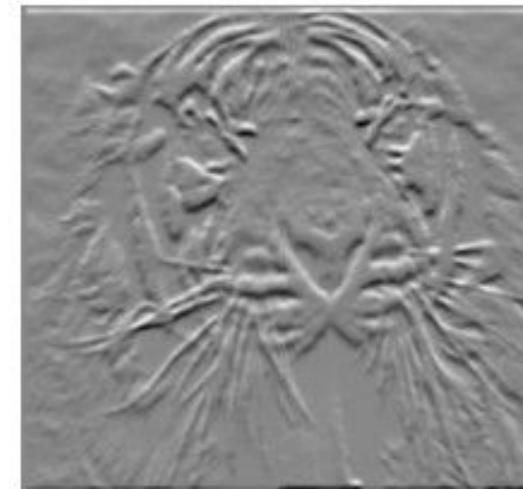
$$\frac{\partial f(x, y)}{\partial x}$$

-1	1
----	---



$$\frac{\partial f(x, y)}{\partial y}$$

-1
1



Which shows changes with respect to x?

Finite difference filters

- Other approximations of derivative filters exist:

Prewitt

-1	0	1
-1	0	1
-1	0	1

Sobel

-1	0	1
-2	0	2
-1	0	1

Roberts

0	1
-1	0

1	1	1
0	0	0
-1	-1	-1

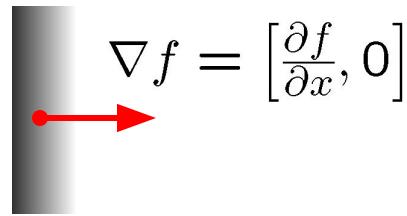
1	2	1
0	0	0
-1	-2	-1

1	0
0	-1

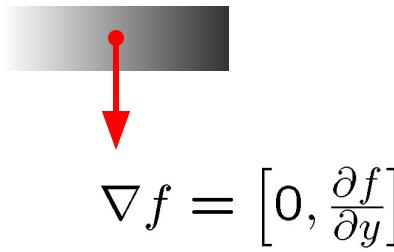
Image gradient

- The gradient of an image:

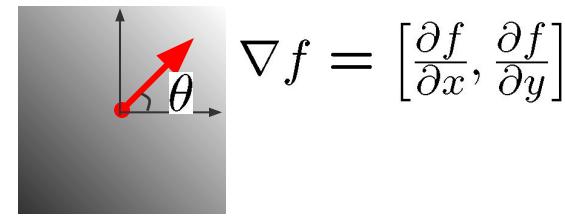
$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$



$$\nabla f = \left[\frac{\partial f}{\partial x}, 0 \right]$$



$$\nabla f = \left[0, \frac{\partial f}{\partial y} \right]$$

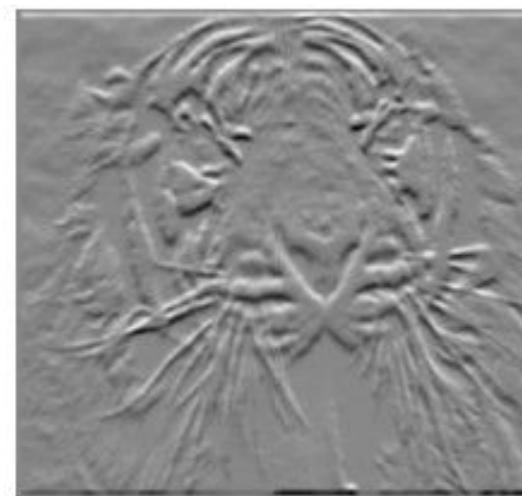
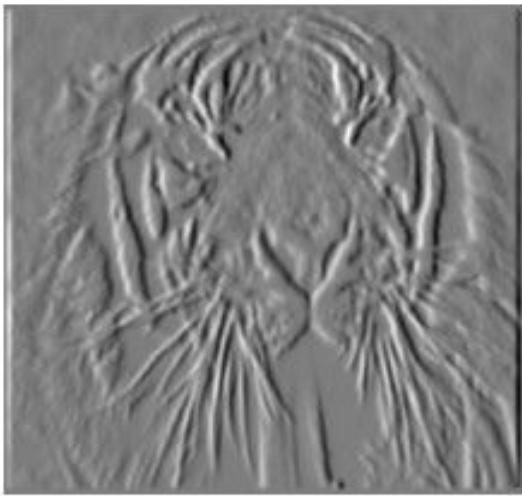


$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

- Gradient points in the direction of most rapid increase in intensity
 - How does this direction relate to the direction of the edge?
- The gradient direction is given by $\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$
- The edge strength is given by the gradient magnitude

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

Image Gradient



$$\frac{\partial f(x, y)}{\partial x}$$

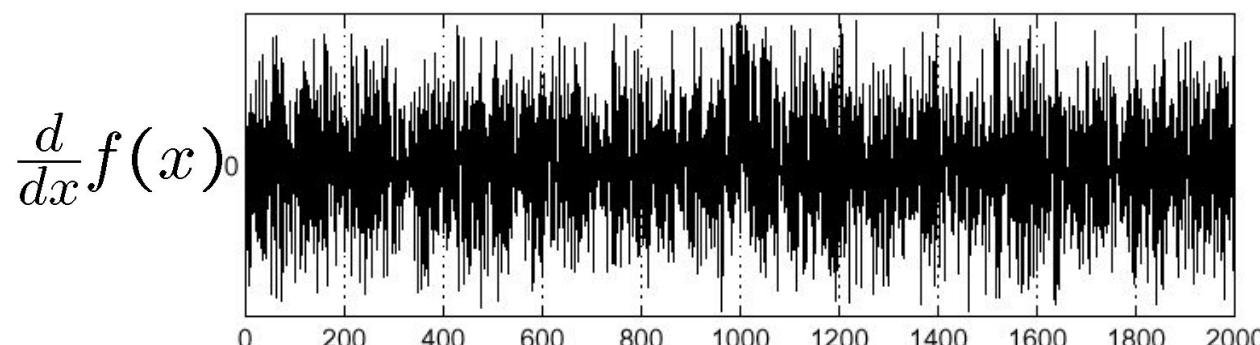
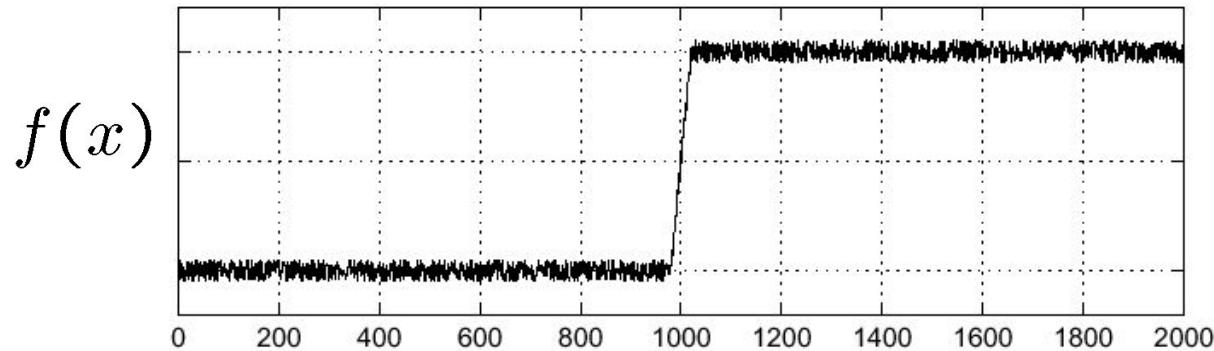
$$\frac{\partial f(x, y)}{\partial y}$$

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$



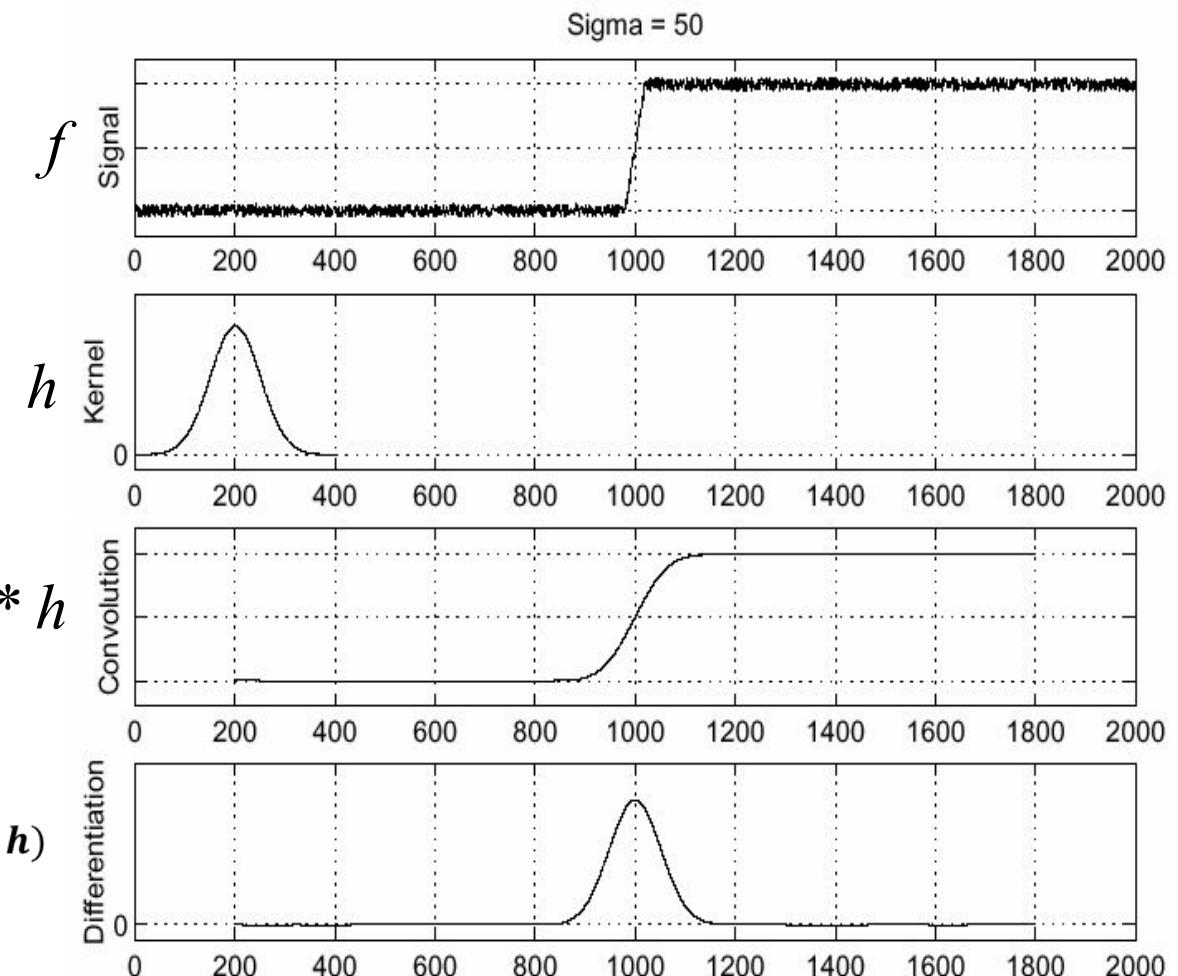
Effects of noise

- Consider a single row the image (shown below)
- How can we find the edge?



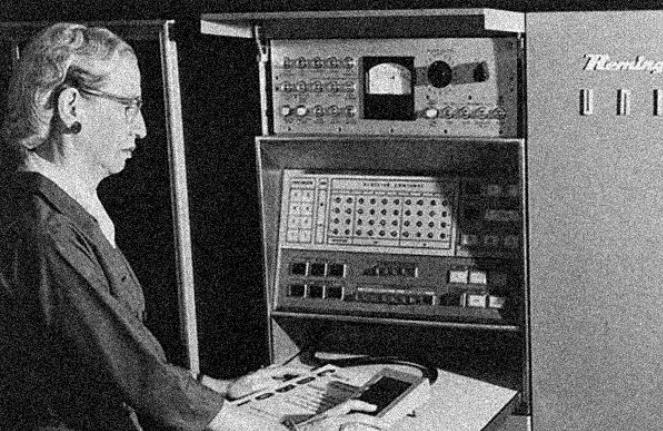
Solution: smooth first

- Look for peak in $\frac{d}{dx}(f * h)$ to find the edge

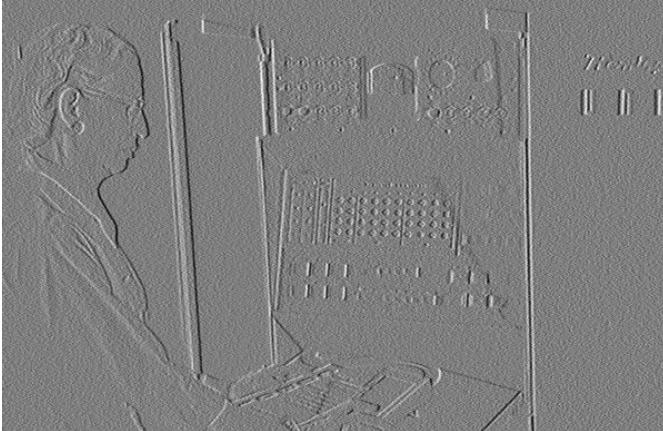


Noise in 2D

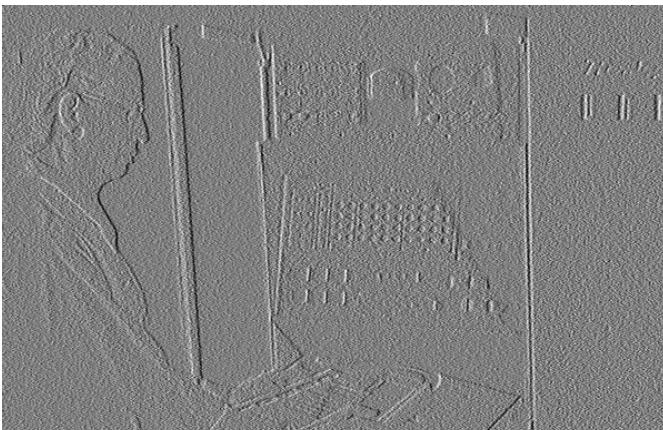
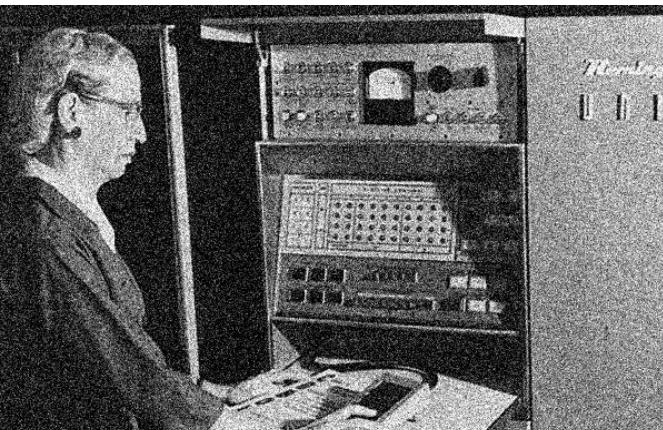
Noisy Input



$\text{dx via } [-1,1]$



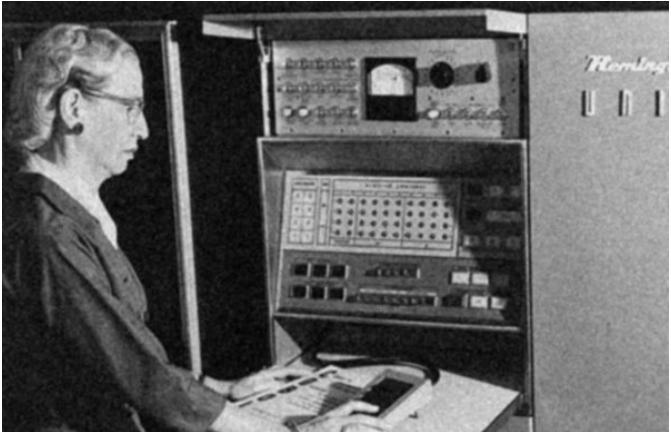
Zoom



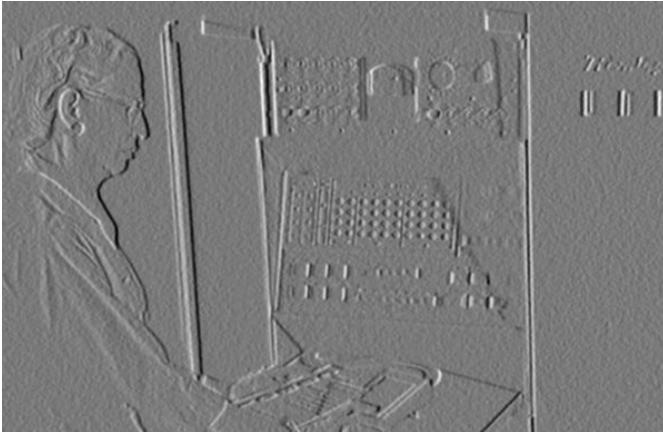
Source: D. Fouhey

Noise + Smoothing

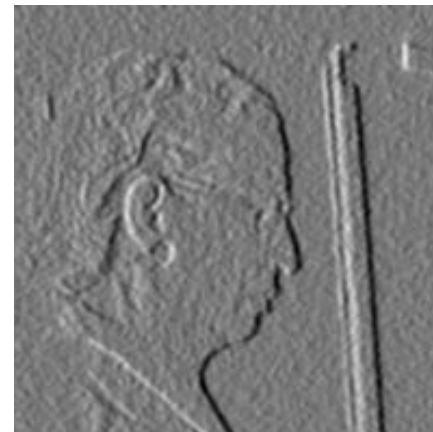
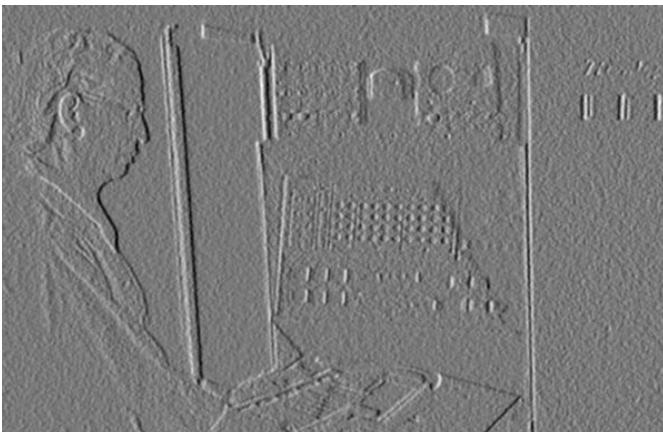
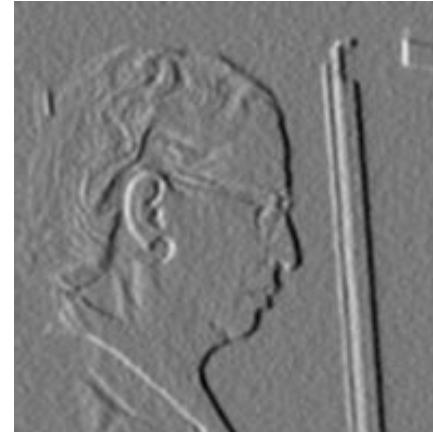
Smoothed Input



dx via $[-1,1]$



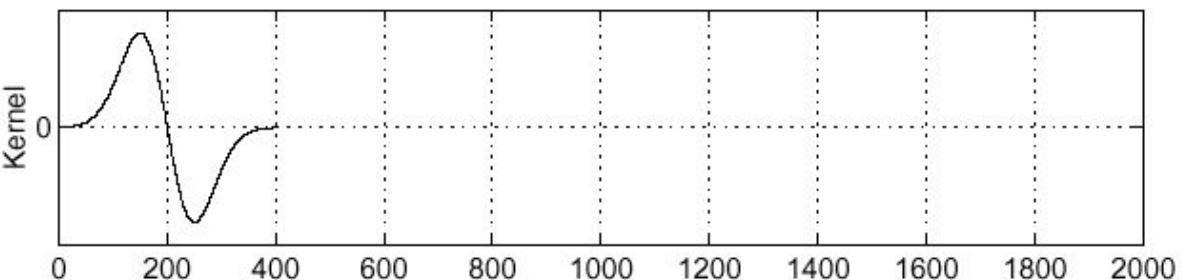
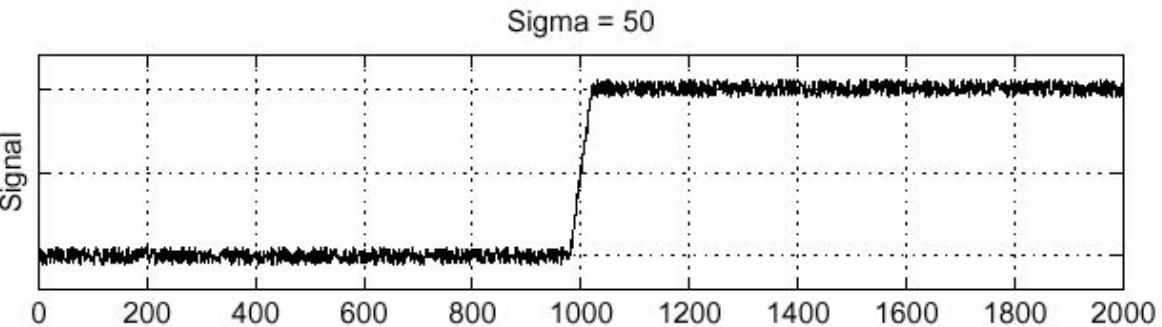
Zoom



Source: D. Fouhey

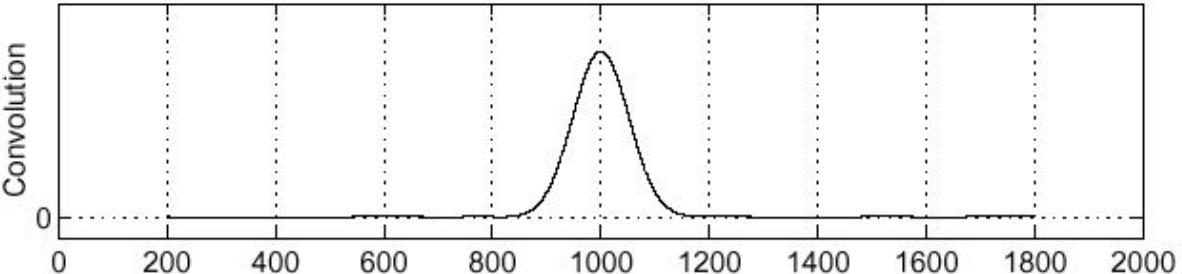
Derivative theorem of convolution

- This saves us one operation: $\frac{\partial}{\partial x}(h \star f) = (\frac{\partial}{\partial x}h) \star f$

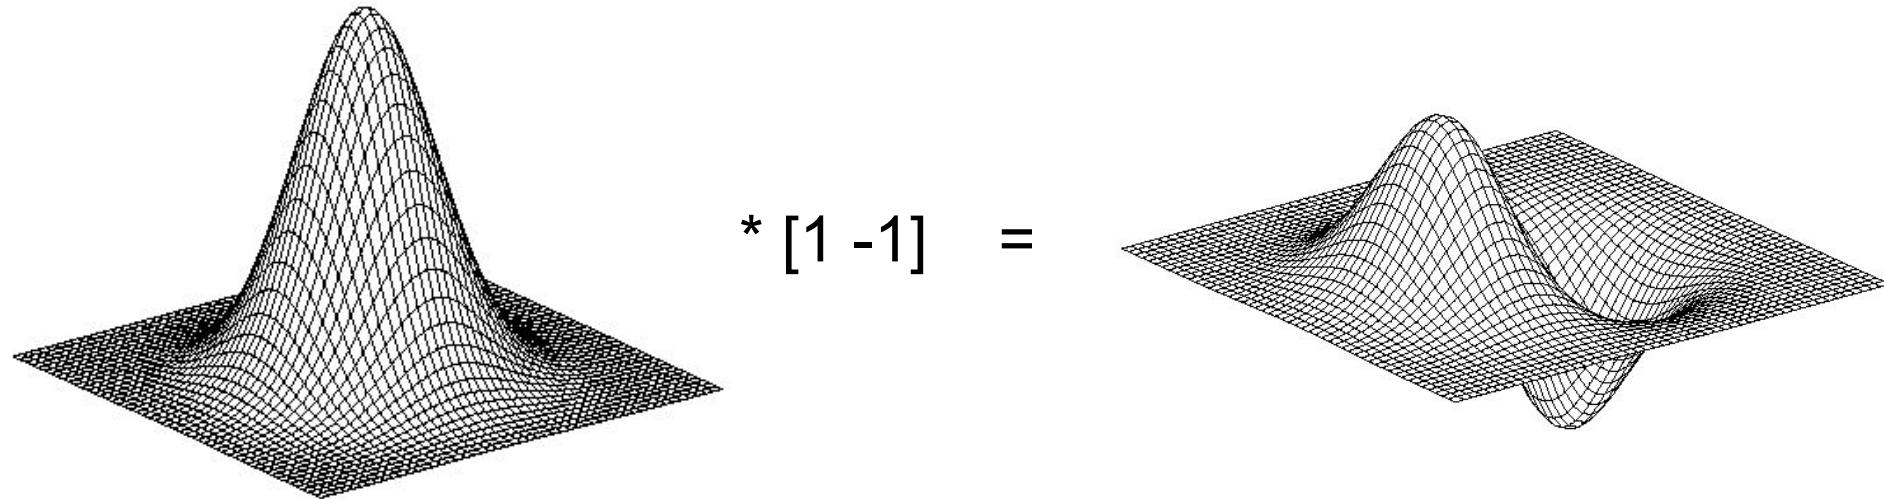


$$\frac{\partial}{\partial x}h$$

$$(\frac{\partial}{\partial x}h) \star f$$

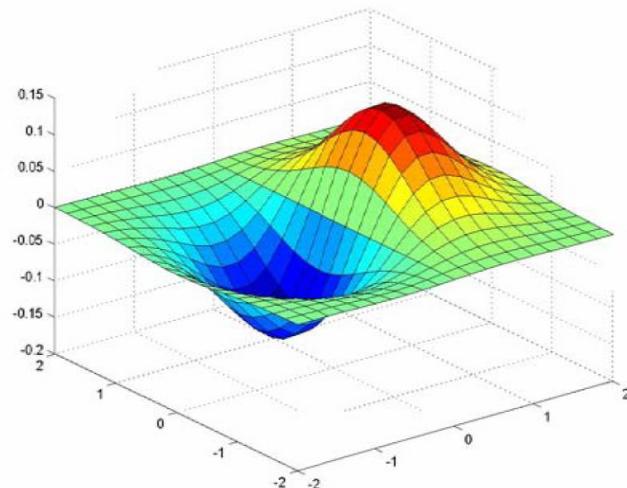


Derivative of Gaussian filter

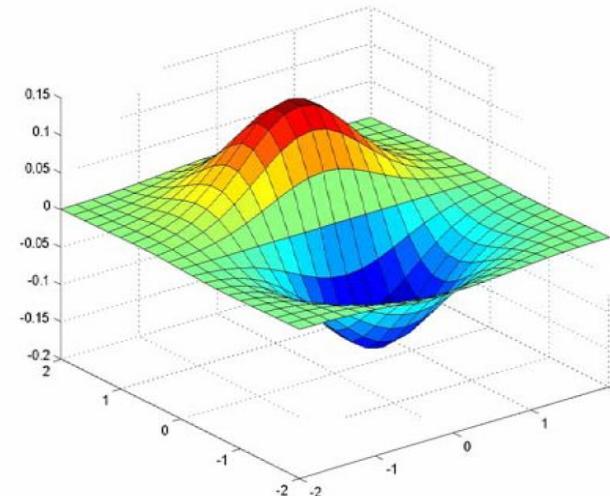


Derivative of Gaussian filter

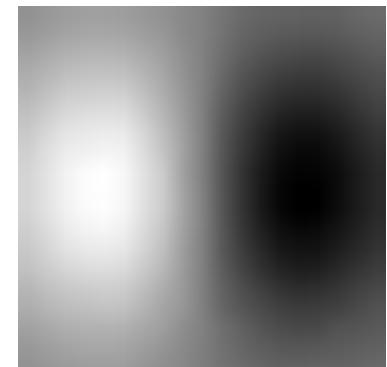
- Which one finds horizontal/vertical edges?



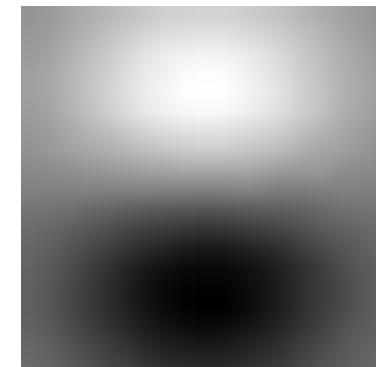
x-direction



y-direction



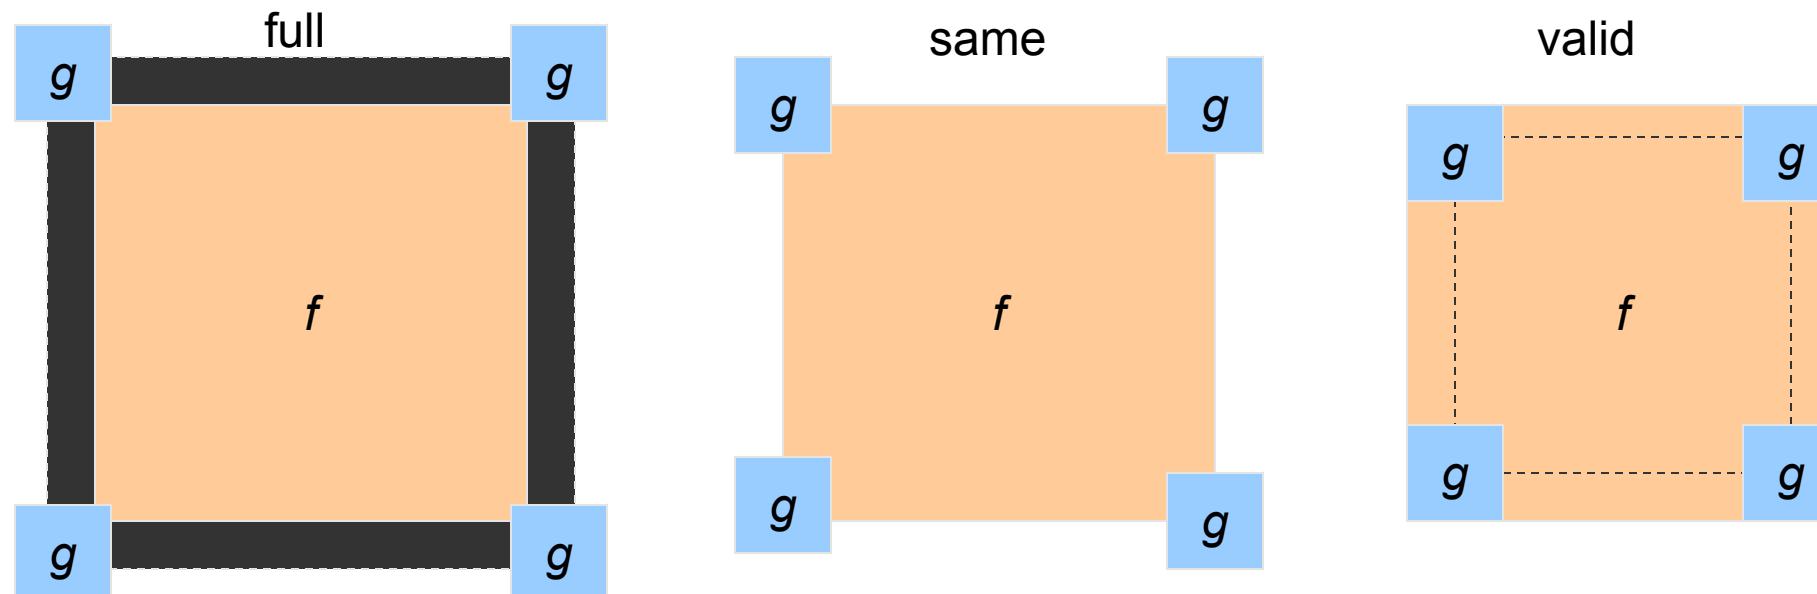
Vertical



Horizontal

Practical matters

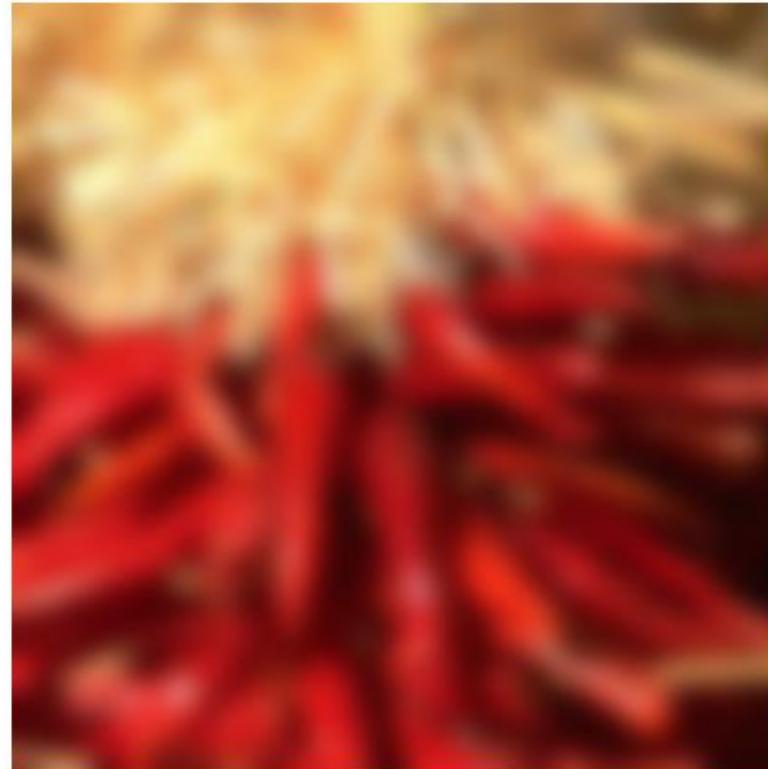
- What is the size of the output?
- Python: `convolve2d(f, g, shape)`
 - *shape = 'full'*: output size is sum of sizes of f and g
 - *shape = 'same'*: output size is same as f
 - *shape = 'valid'*: output size is difference of sizes of f and g



Practical matters

□ What about near the edge?

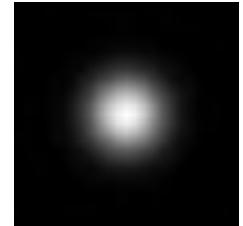
- the filter window falls off the edge of the image
- need to extrapolate
- methods:
 - clip filter (black)
 - wrap around
 - copy edge
 - reflect across edge



Summary

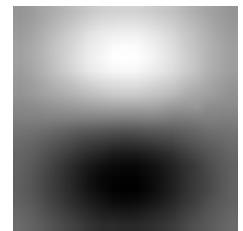
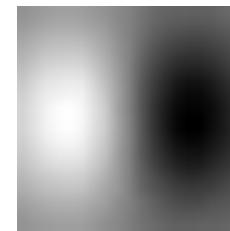
□ Smoothing filters

- Gaussian: remove “high-frequency” components; “low-pass” filter
- Can the values of a smoothing filter be negative?
- What should the values sum to?
 - One: constant regions are not affected by the filter



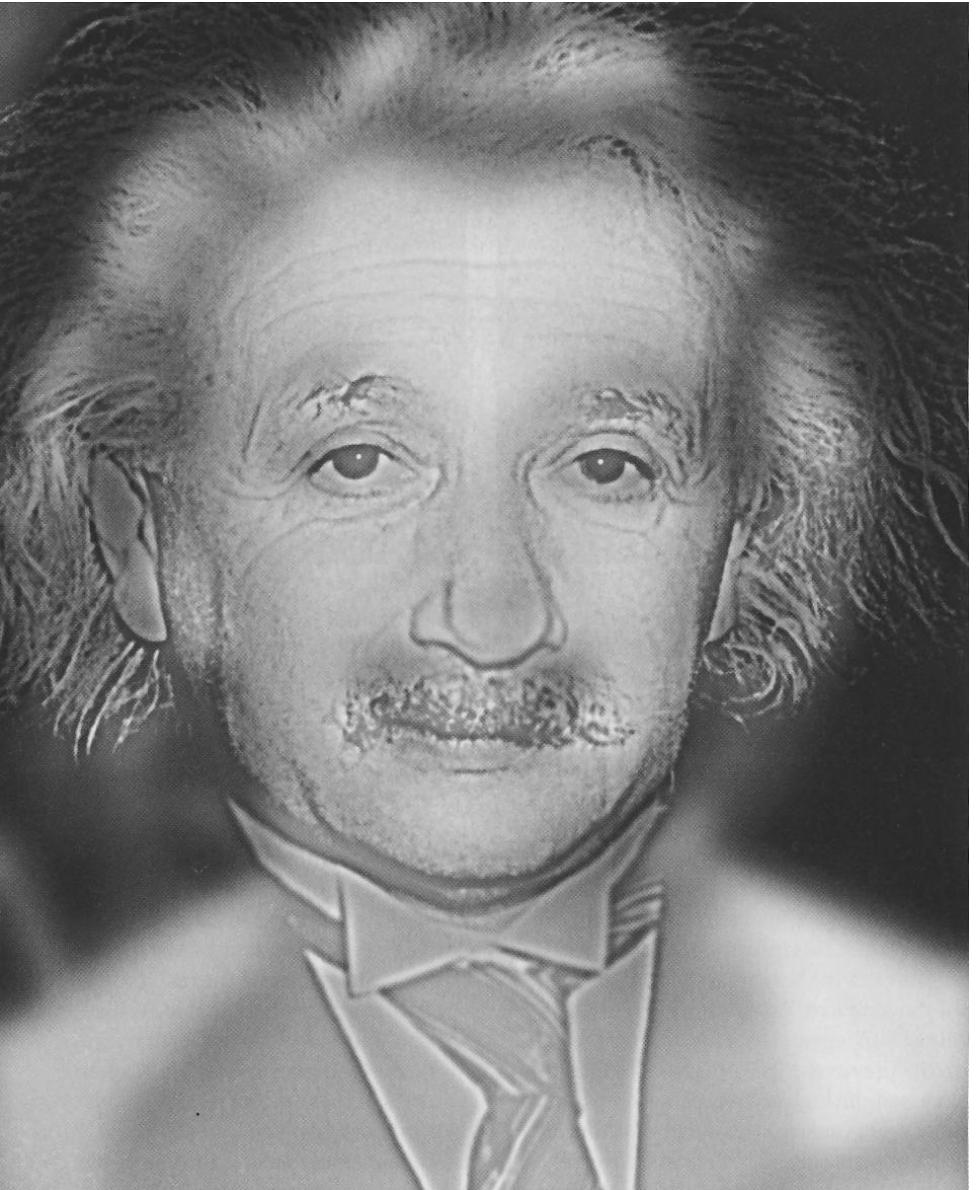
□ Derivative filters

- Derivatives of Gaussian
- Can the values of a derivative filter be negative?
- What should the values sum to?
 - Zero: no response in constant regions

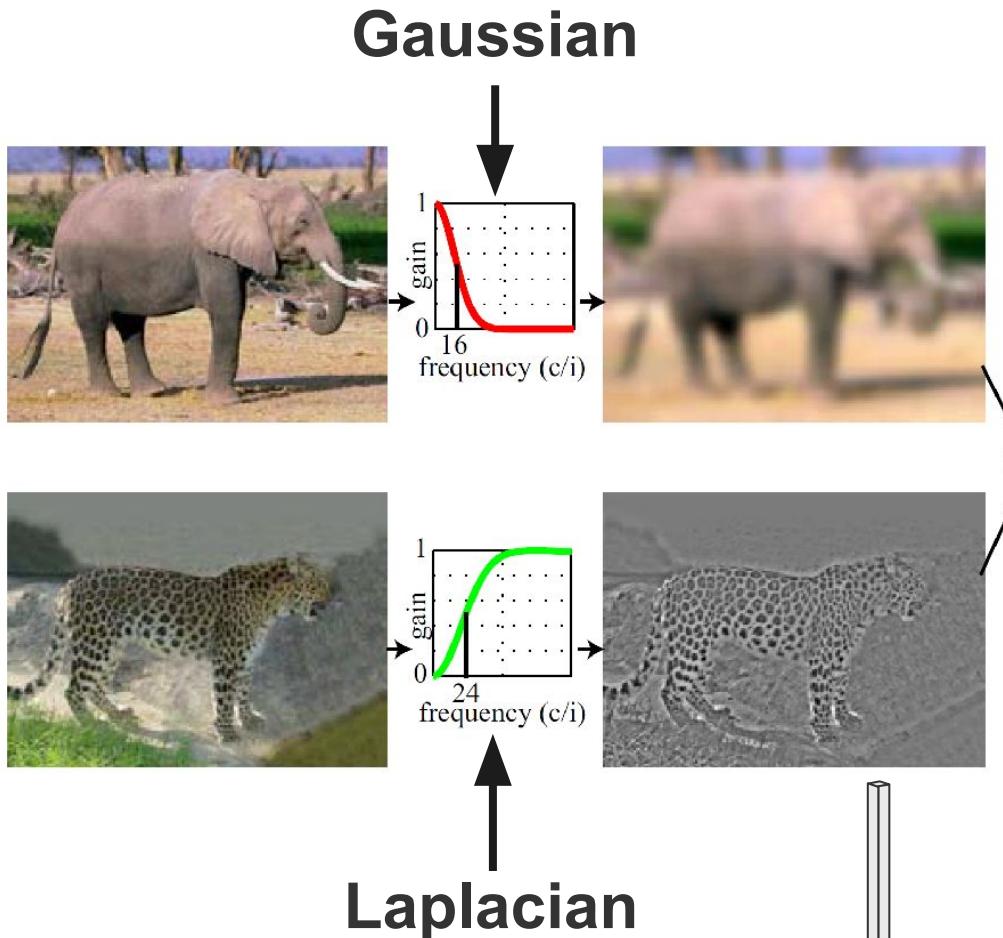


A cool application

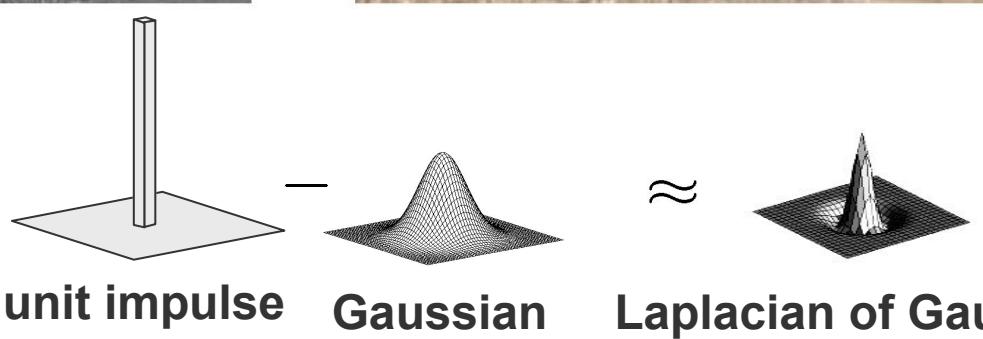
Hybrid images



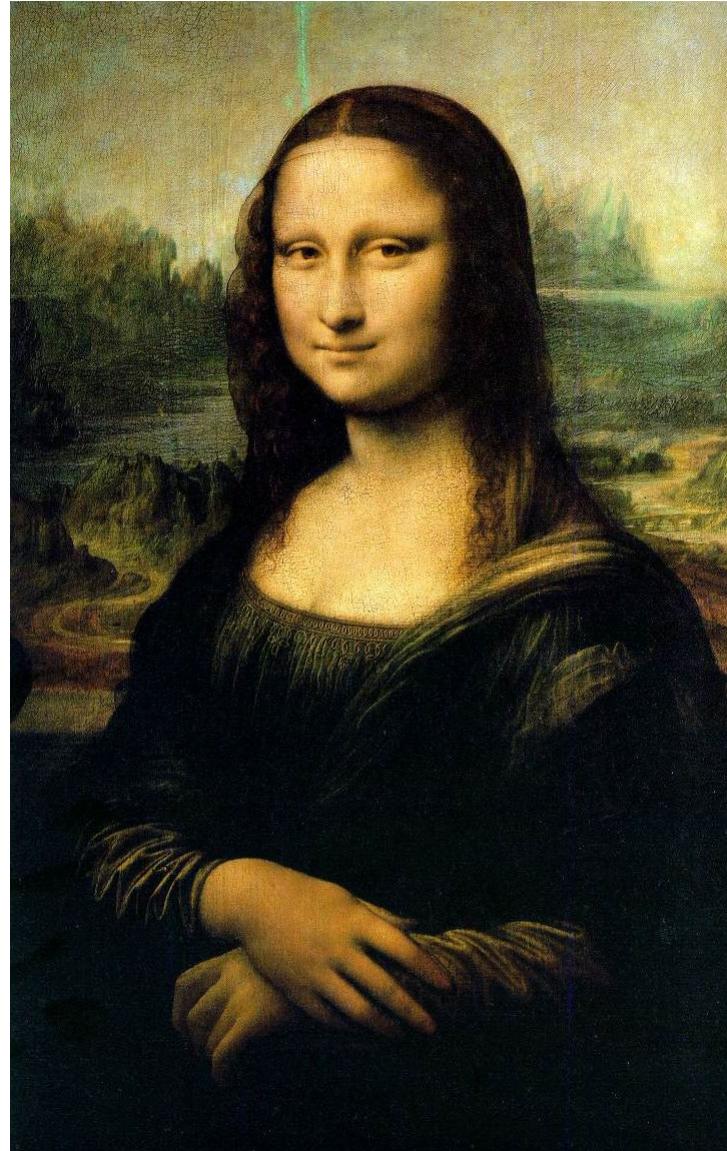
Hybrid Images



A. Oliva, A. Torralba, P.G. Schyns,
“Hybrid Images,” SIGGRAPH 2006



Da Vinci and Peripheral Vision



Da Vinci and Peripheral Vision

□ Leonardo playing with peripheral vision

